

# PetPeople





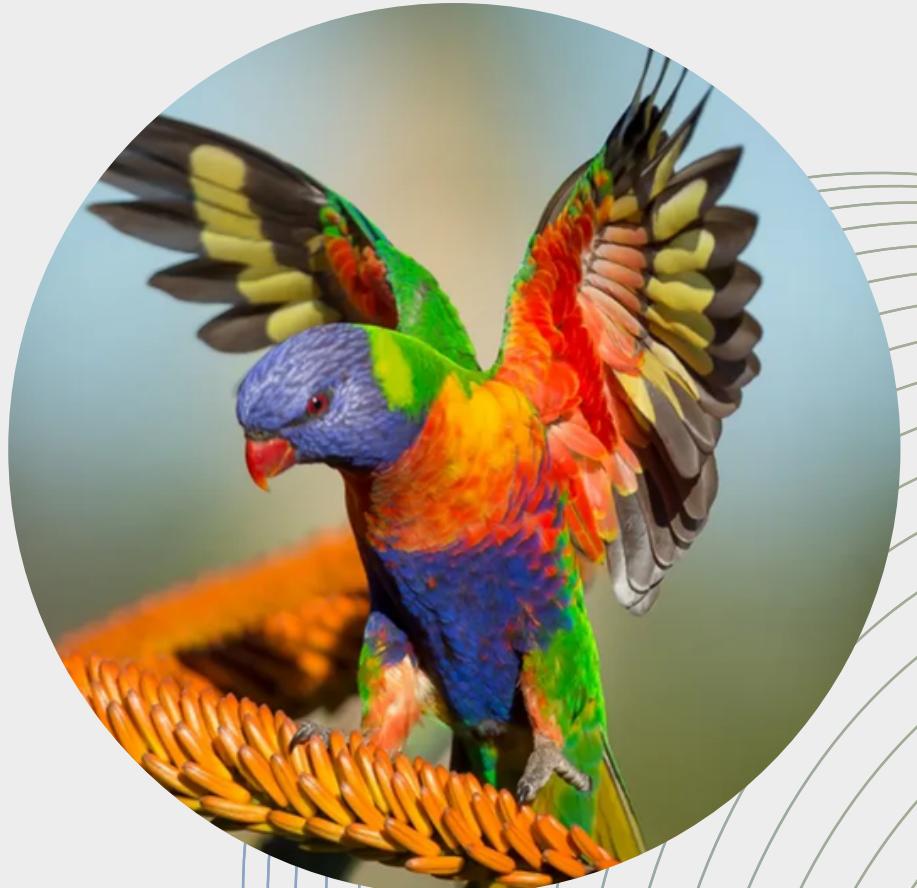
# PetPeople

Social network

MarketPlace

Information Hub

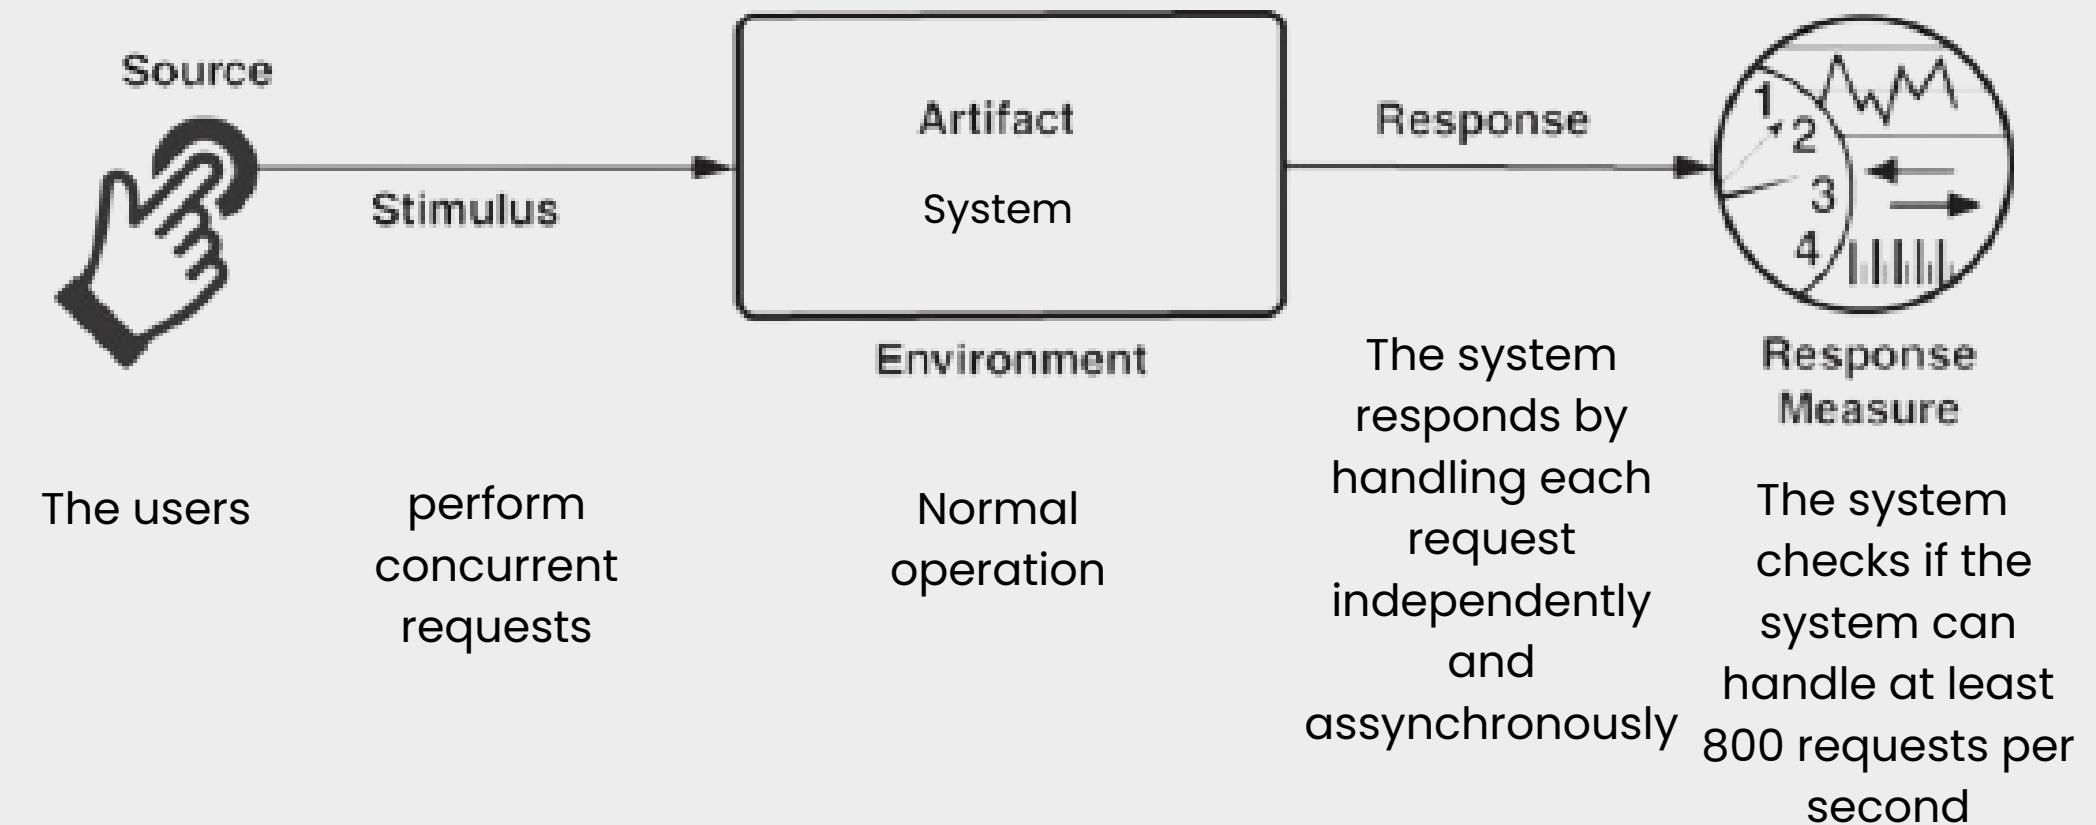
For your pet



# Framework



- Number of concurrent requests (Performance)

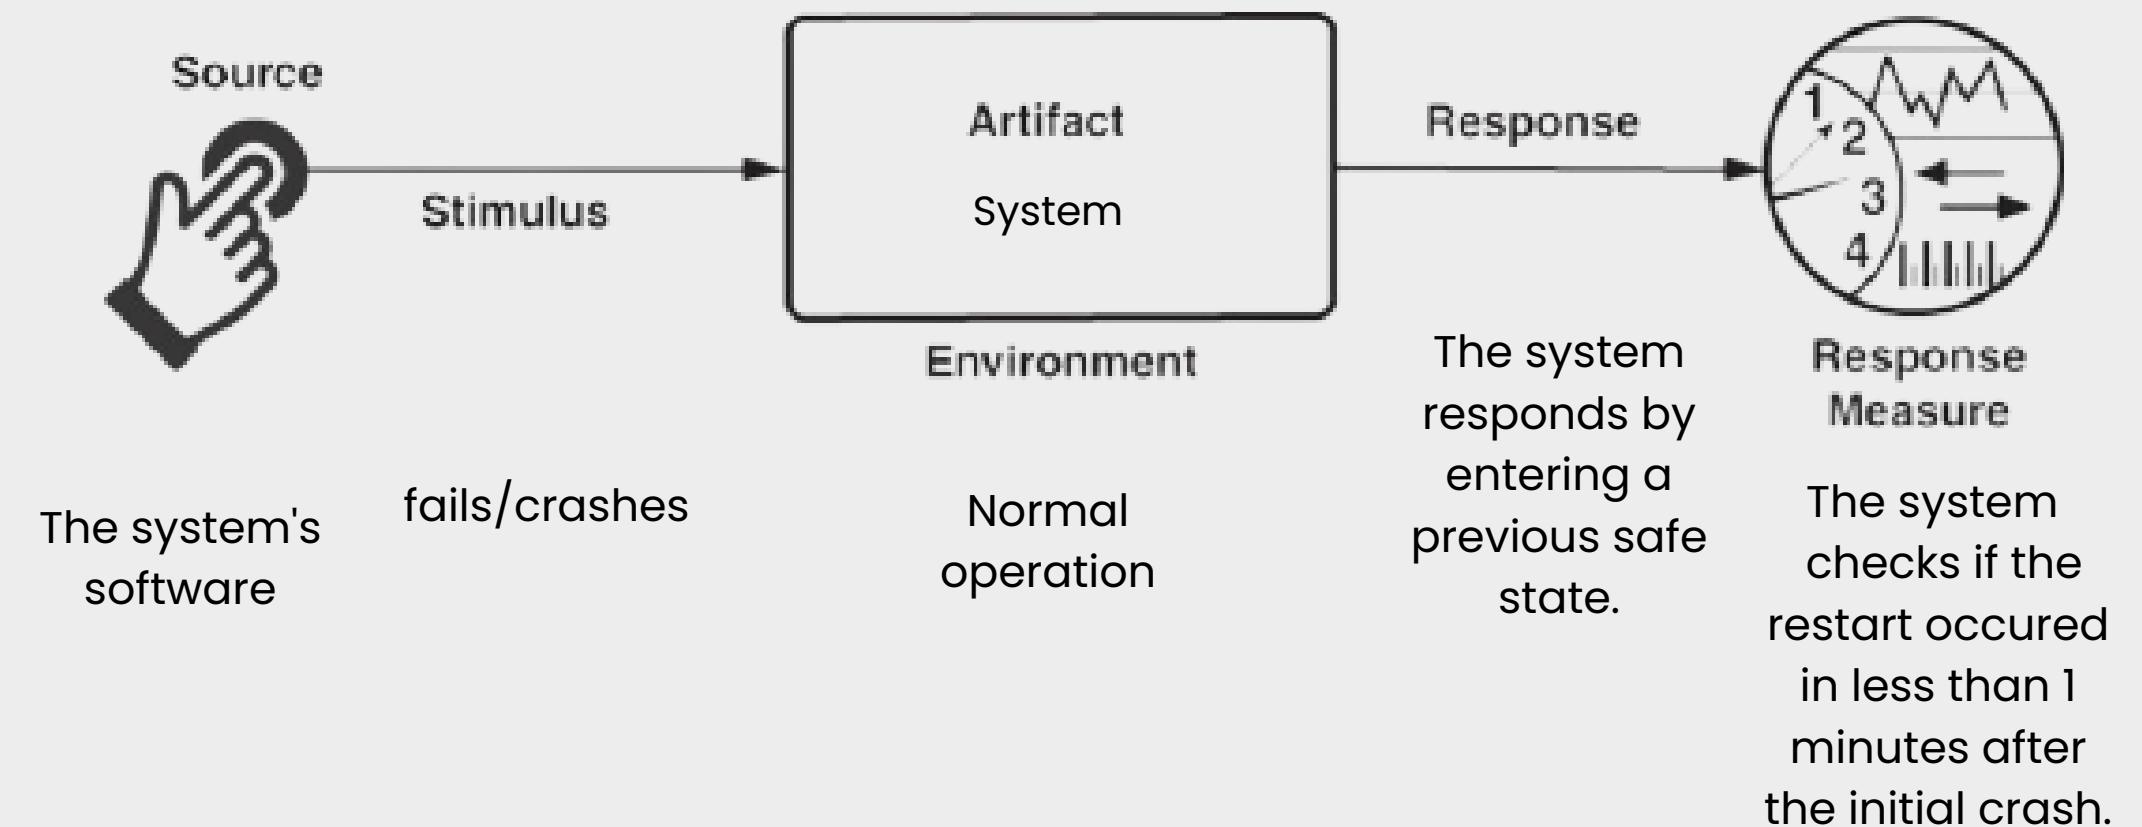


- **Express** is known for its simplicity and speed, making it an excellent choice for building lightweight, high-performance applications.
- **Spring** offers a broad range of features and modules, making it suitable for building complex applications that require scalability and high performance.

# Framework



- Up time (Availability)

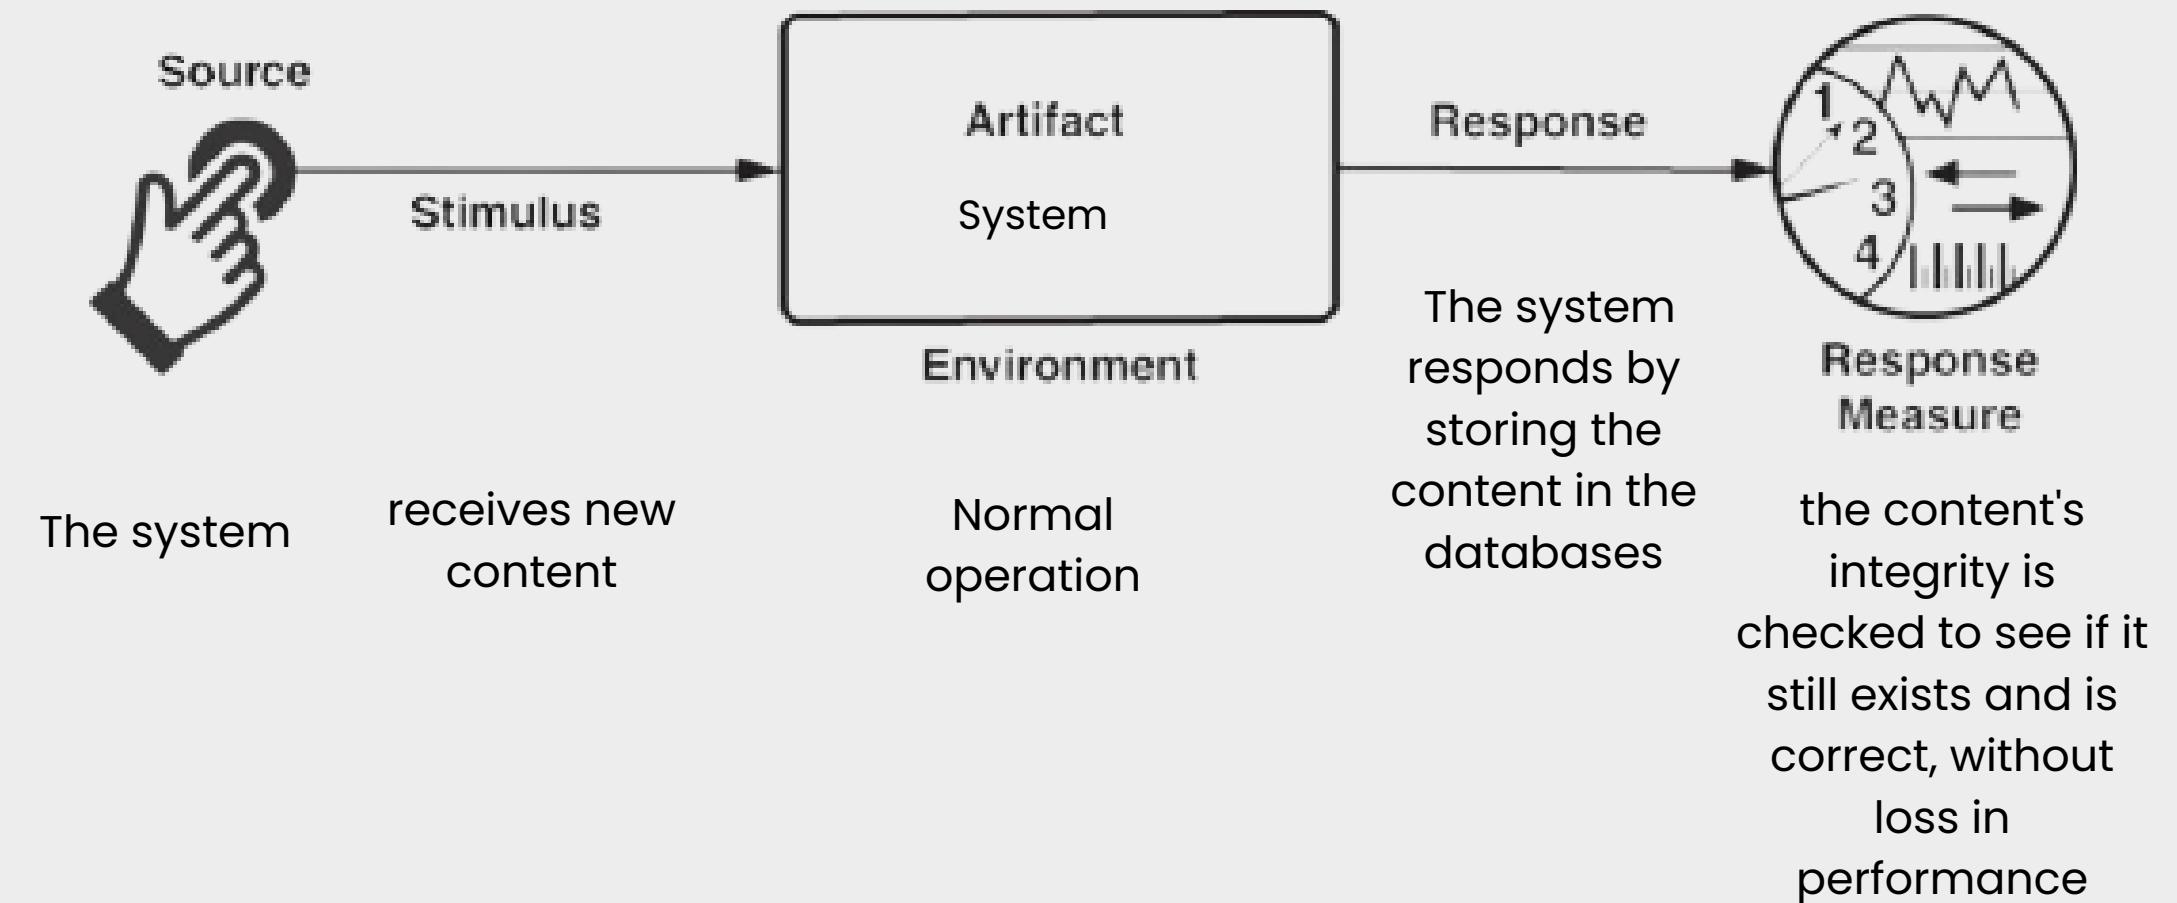


- **Spring**, **Django**, **Laravel**, and **Ruby on Rails** all offer built-in fault tolerance features, while **Express** requires additional third-party libraries and modules. However, the specific fault tolerance features and configuration may differ between frameworks.

# Framework



- Sustain an annual content growth of at least 200% in content (Scalability)

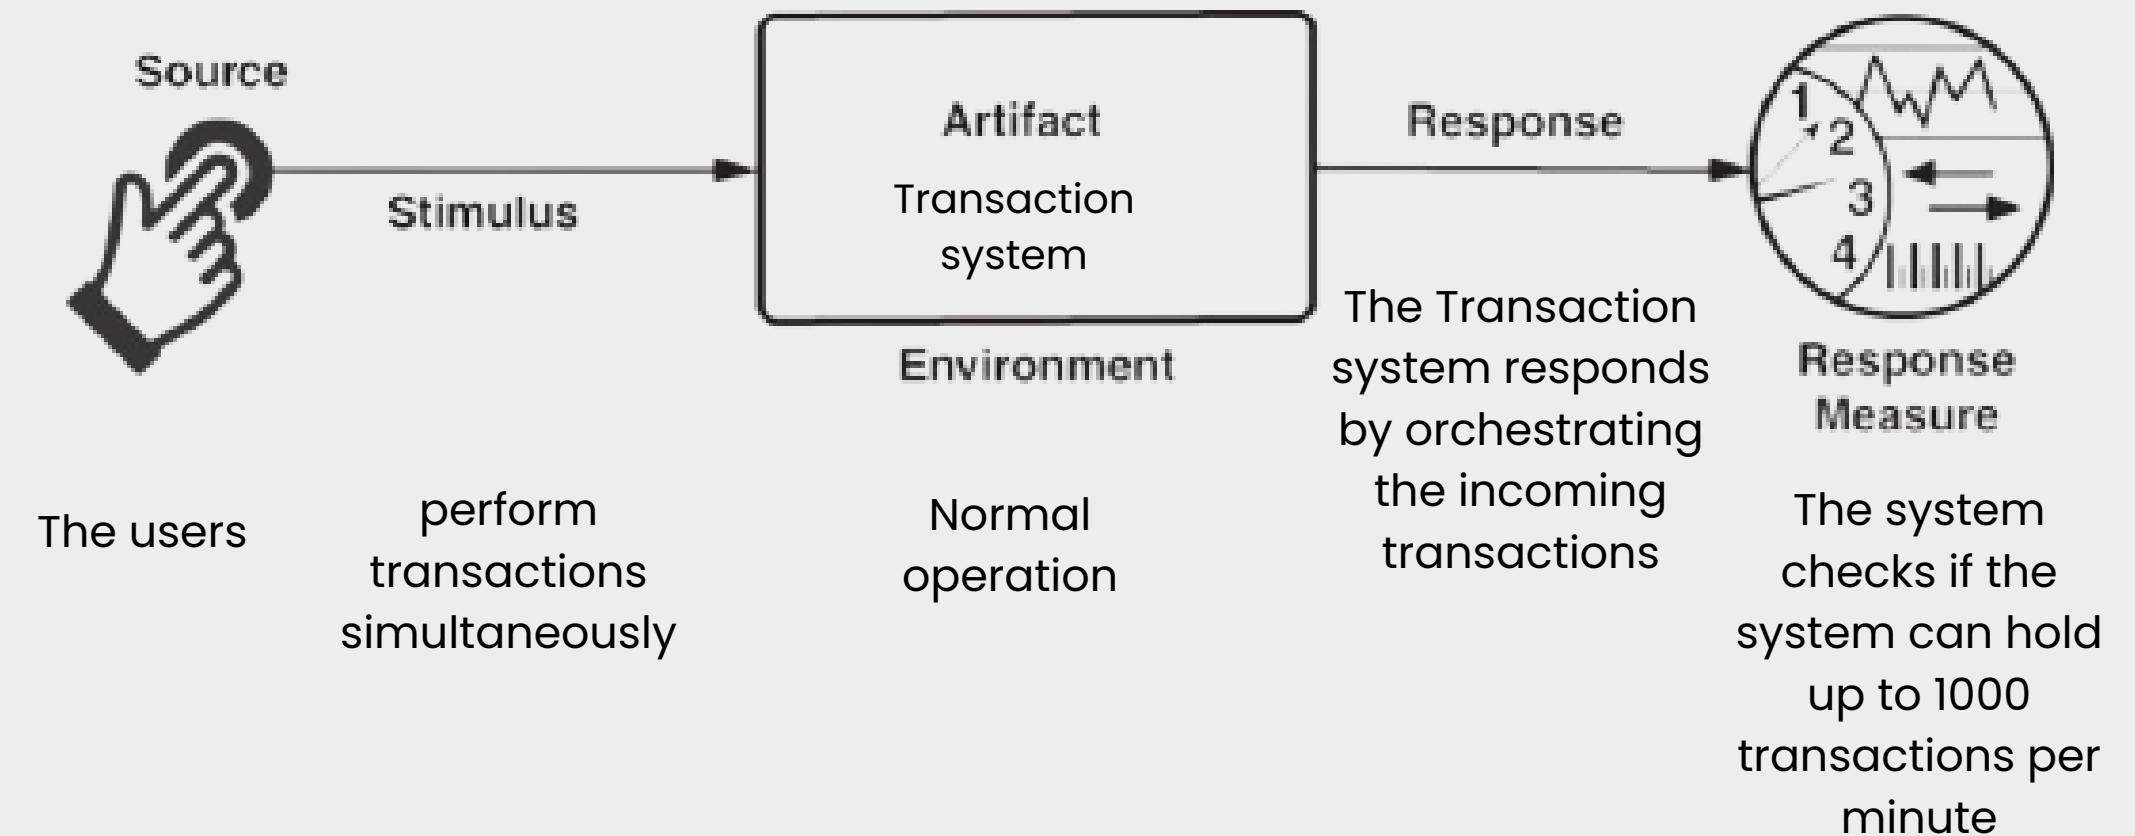


- **Express** is a lightweight and flexible framework that can be easily integrated with different databases and libraries, making it a popular choice for building web applications.
- **Django** has a well-established architecture and comes with a lot of built-in features and plugins, making it easy to develop and deploy web applications.

# Framework



- Number of transactions (Performance)

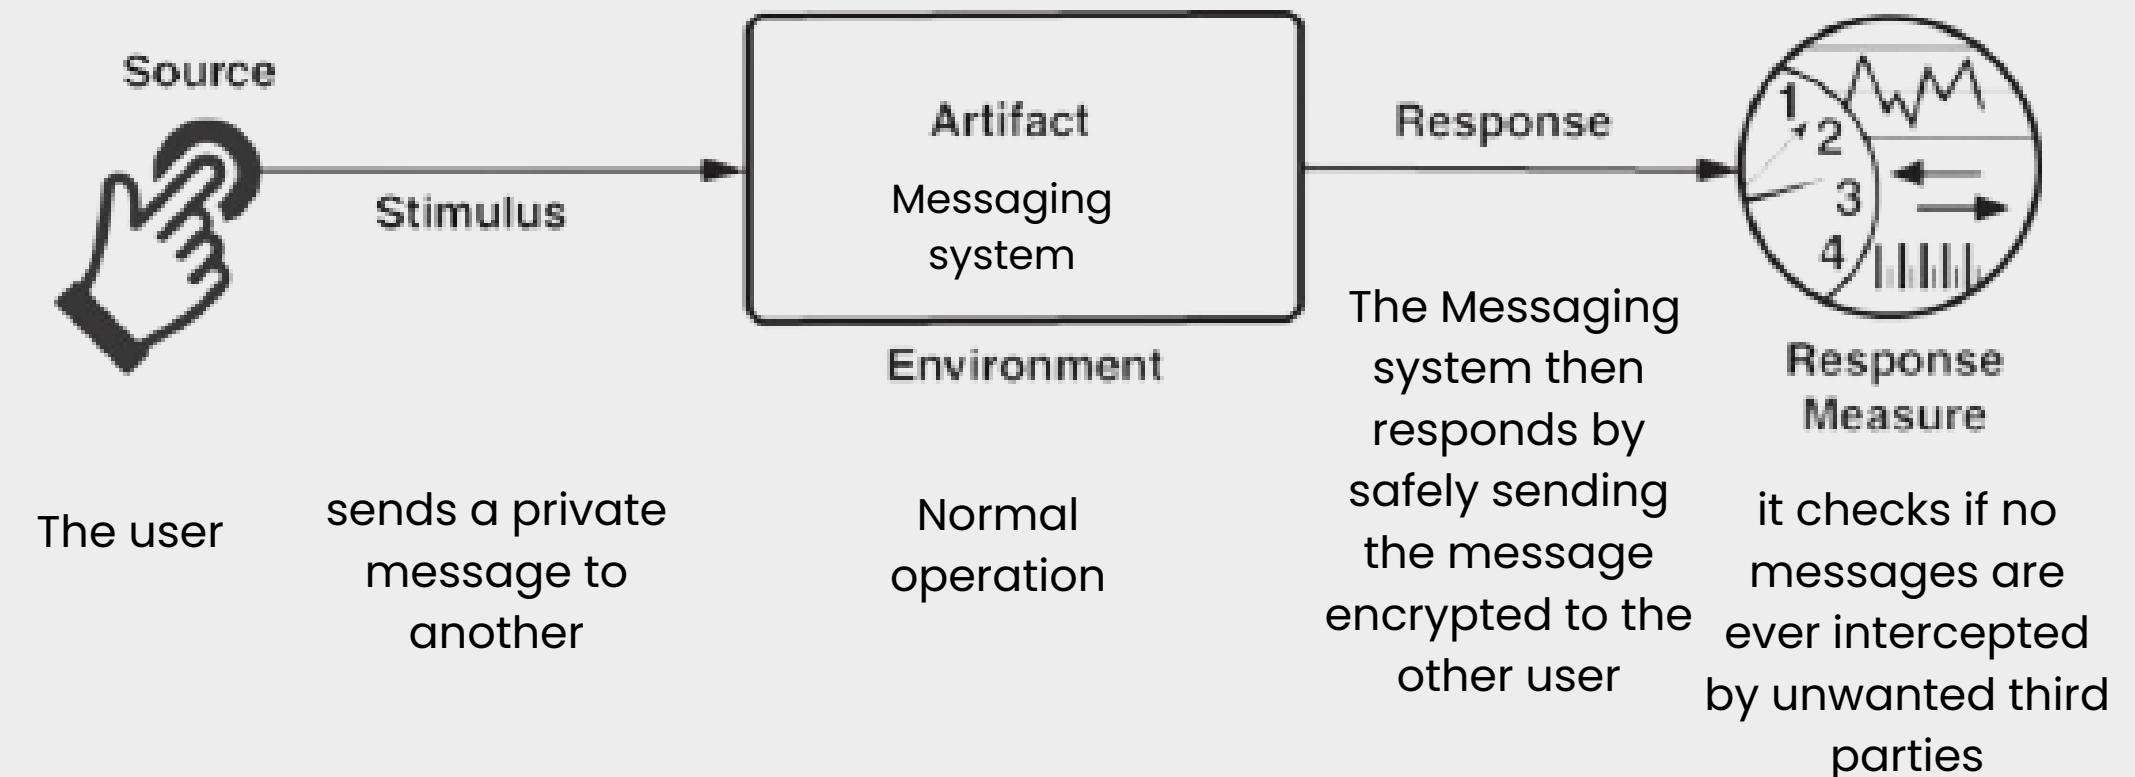


- **Spring** is known for its support for distributed systems and microservices, making it a good choice for building transactional systems that can handle a large volume of transactions.
- **Express**, on the other hand, can handle a high volume of requests concurrently, making it a suitable option for systems that require high concurrency.

# Framework



- Secure messaging (Security)



- Spring and Django** are known for their strong security features, making them a good choice for building secure messaging systems. **Express, Laravel, and Ruby on Rails** can also be used, but you may need to add additional security features to ensure that messages are not intercepted by unwanted third parties.

# Framework



- Low Funds
  - As a Client that is starting a business, I want to save the most I can seeing that I do not have a lot of funds

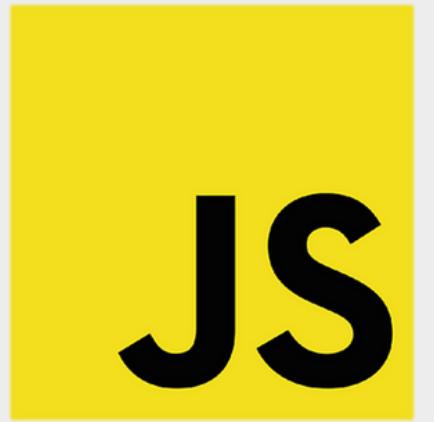
Considering that we are the developers of the platform, for cost savings and a faster development process, the best framework for us would be one with which we are familiar, such as **Express** and **Ruby on Rails**.

# Framework

Decision



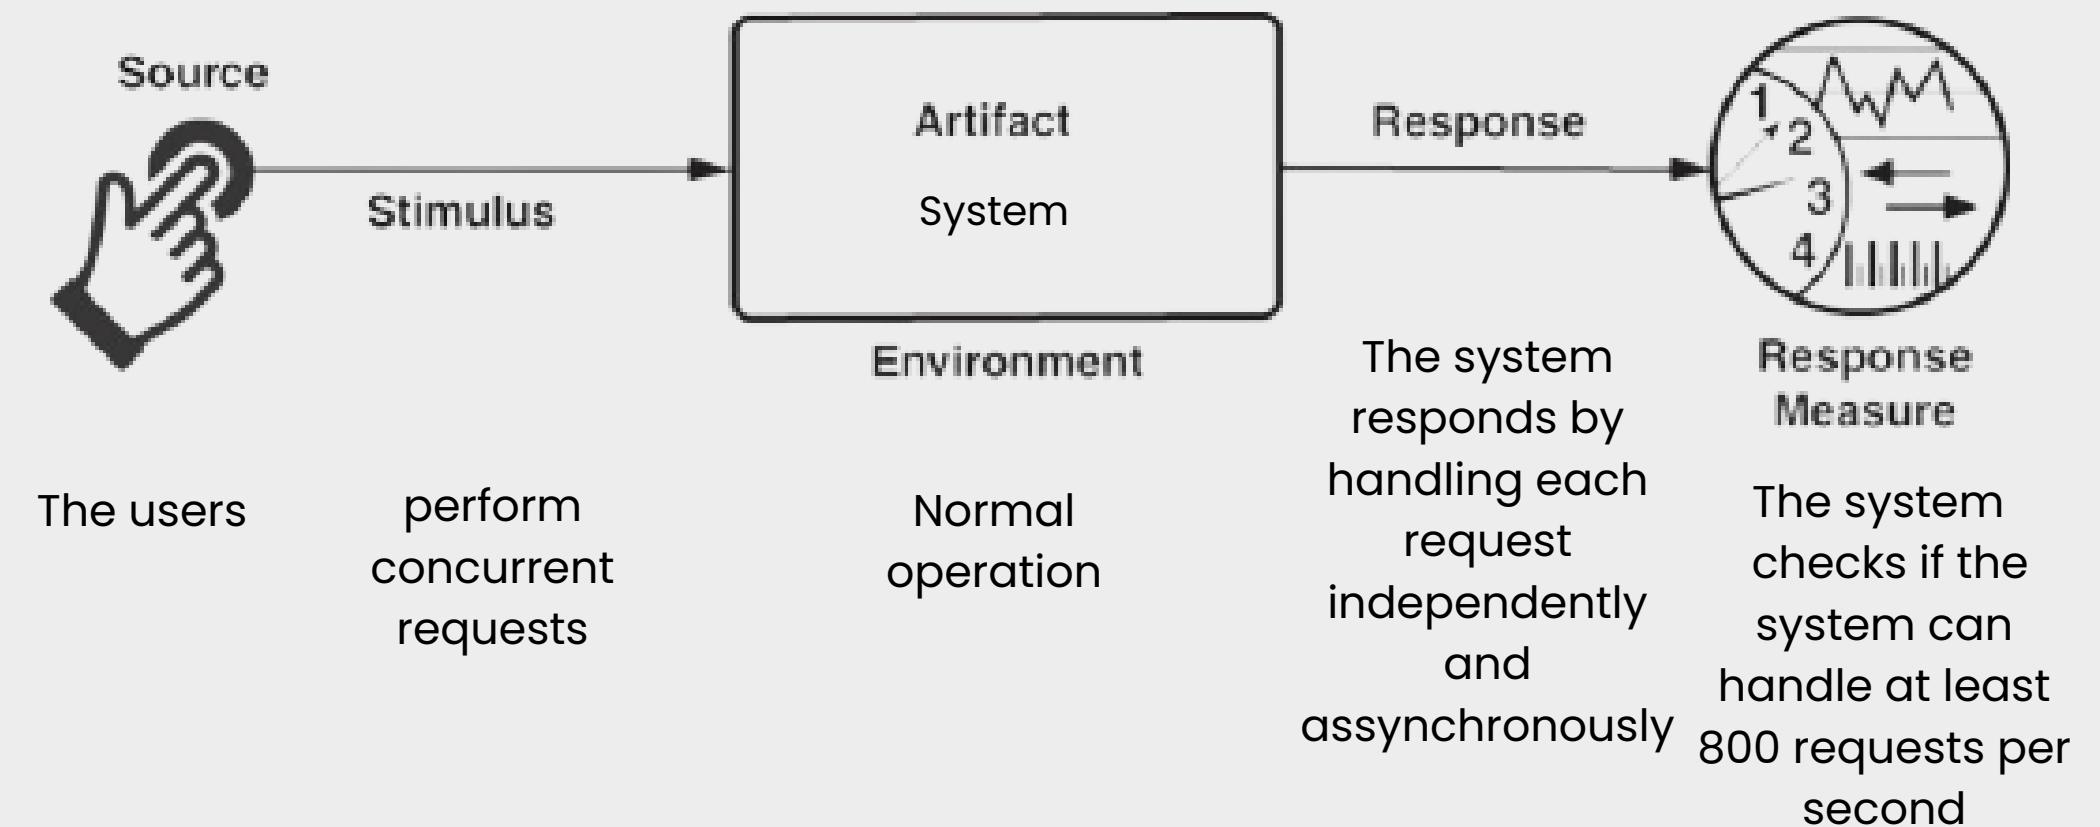
Express



# Database



- Number of concurrent requests (Performance)

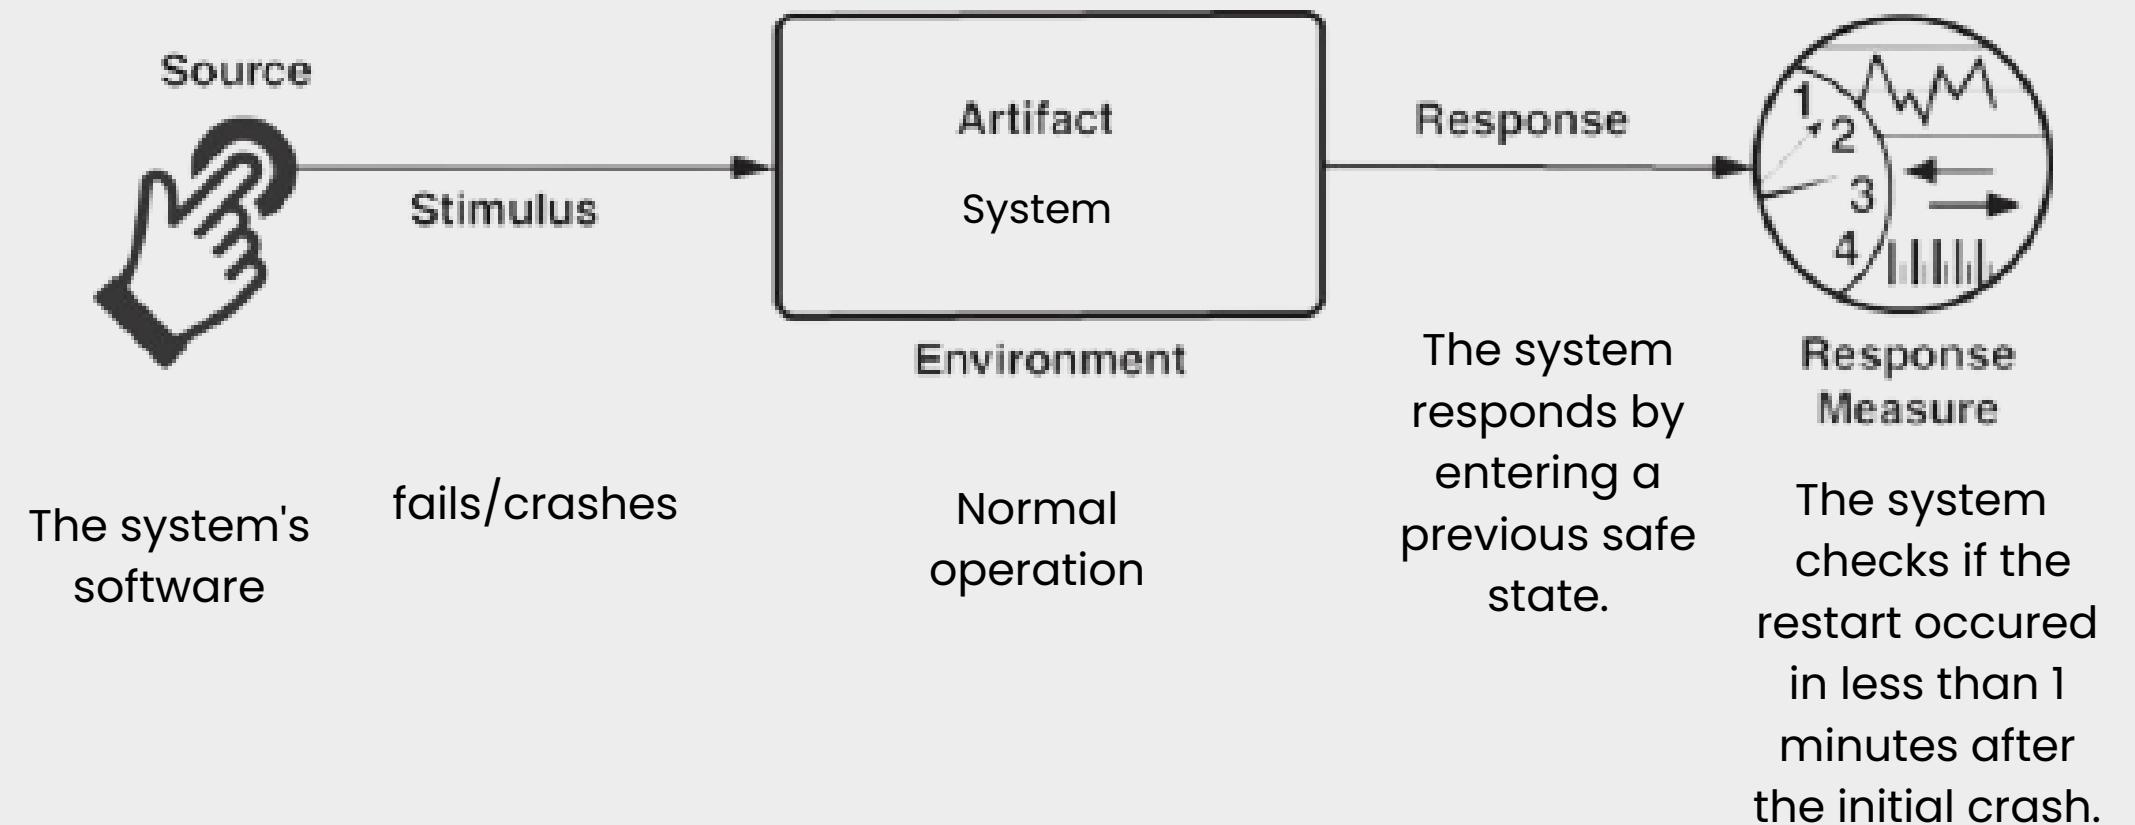


Document-oriented databases are designed for high performance and scalability, and they can handle concurrent requests efficiently. They store data in a flexible, semi-structured format, making them a good fit for web and mobile applications that deal with large volumes of unstructured or semi-structured data. **MongoDB**, in particular, has features like sharding and replication, which allow it to scale horizontally and handle high volumes of traffic.

# Database



- Up time (Availability)

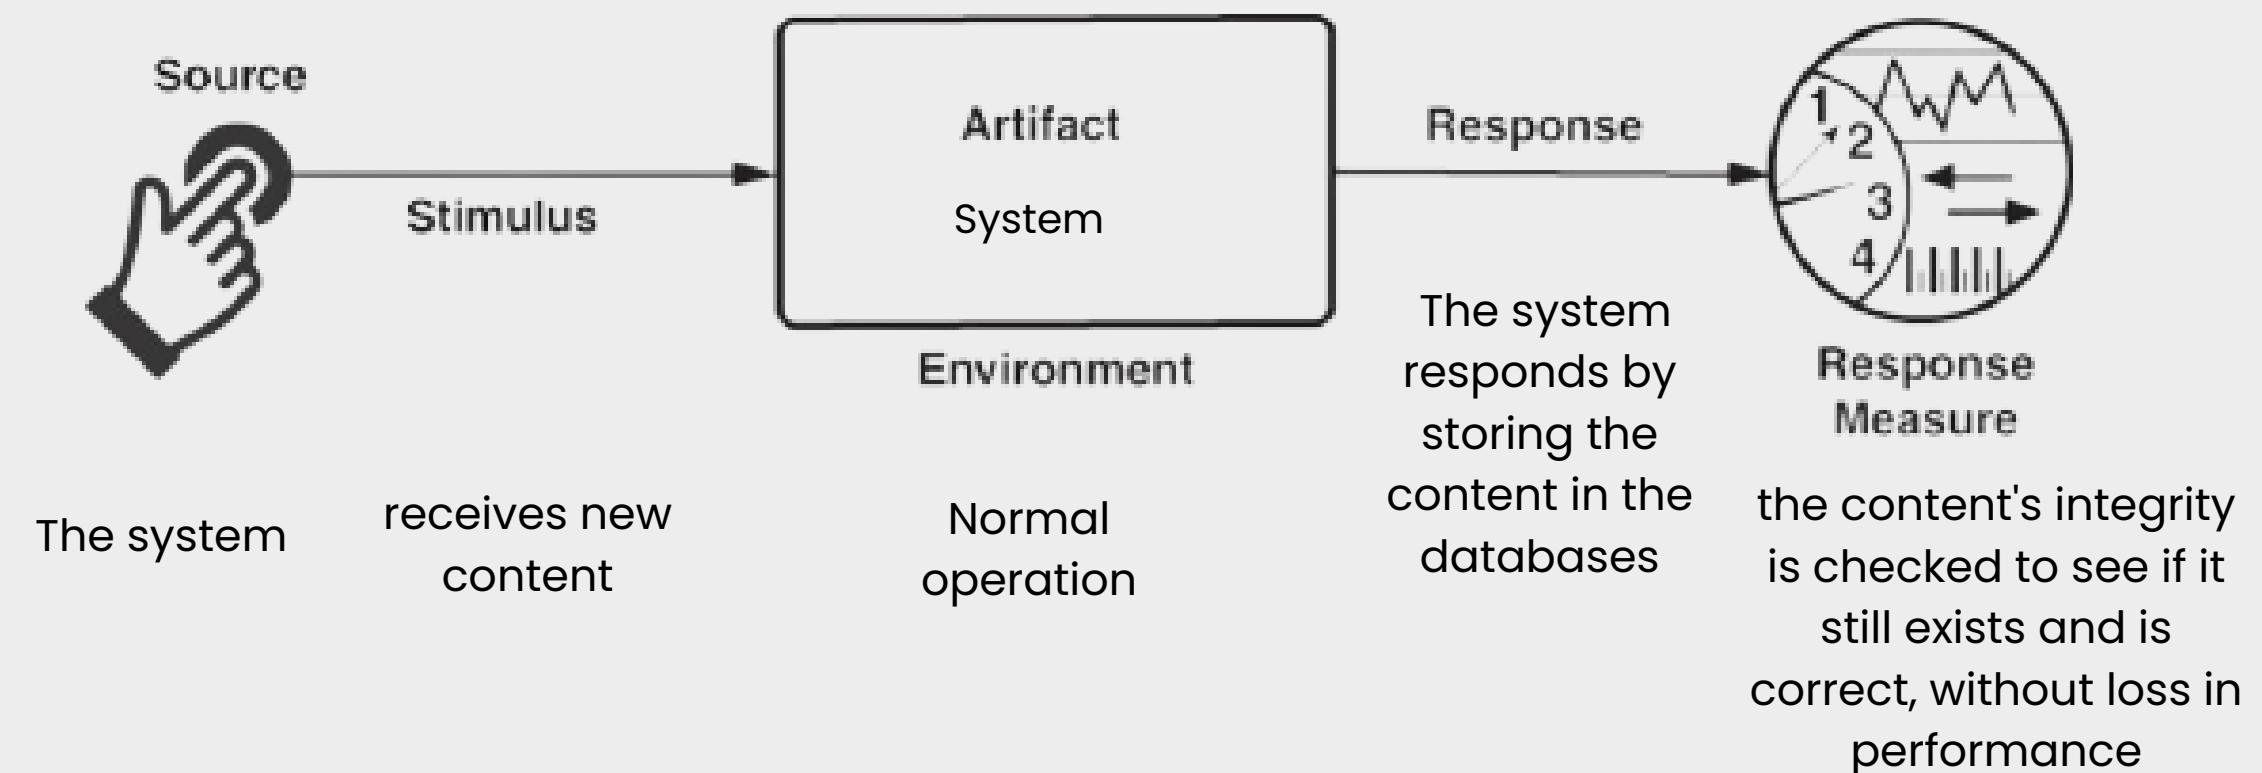


**MySQL**, **PostgreSQL**, and **Oracle** all have relatively mature mechanisms for recovery, such as replication and clustering, which can provide high availability and reduce downtime. **MongoDB** and **Redis** have similar mechanisms, but they are less mature and have some limitations.

# Database



- Sustain an annual content growth of at least 200% in content (Scalability)

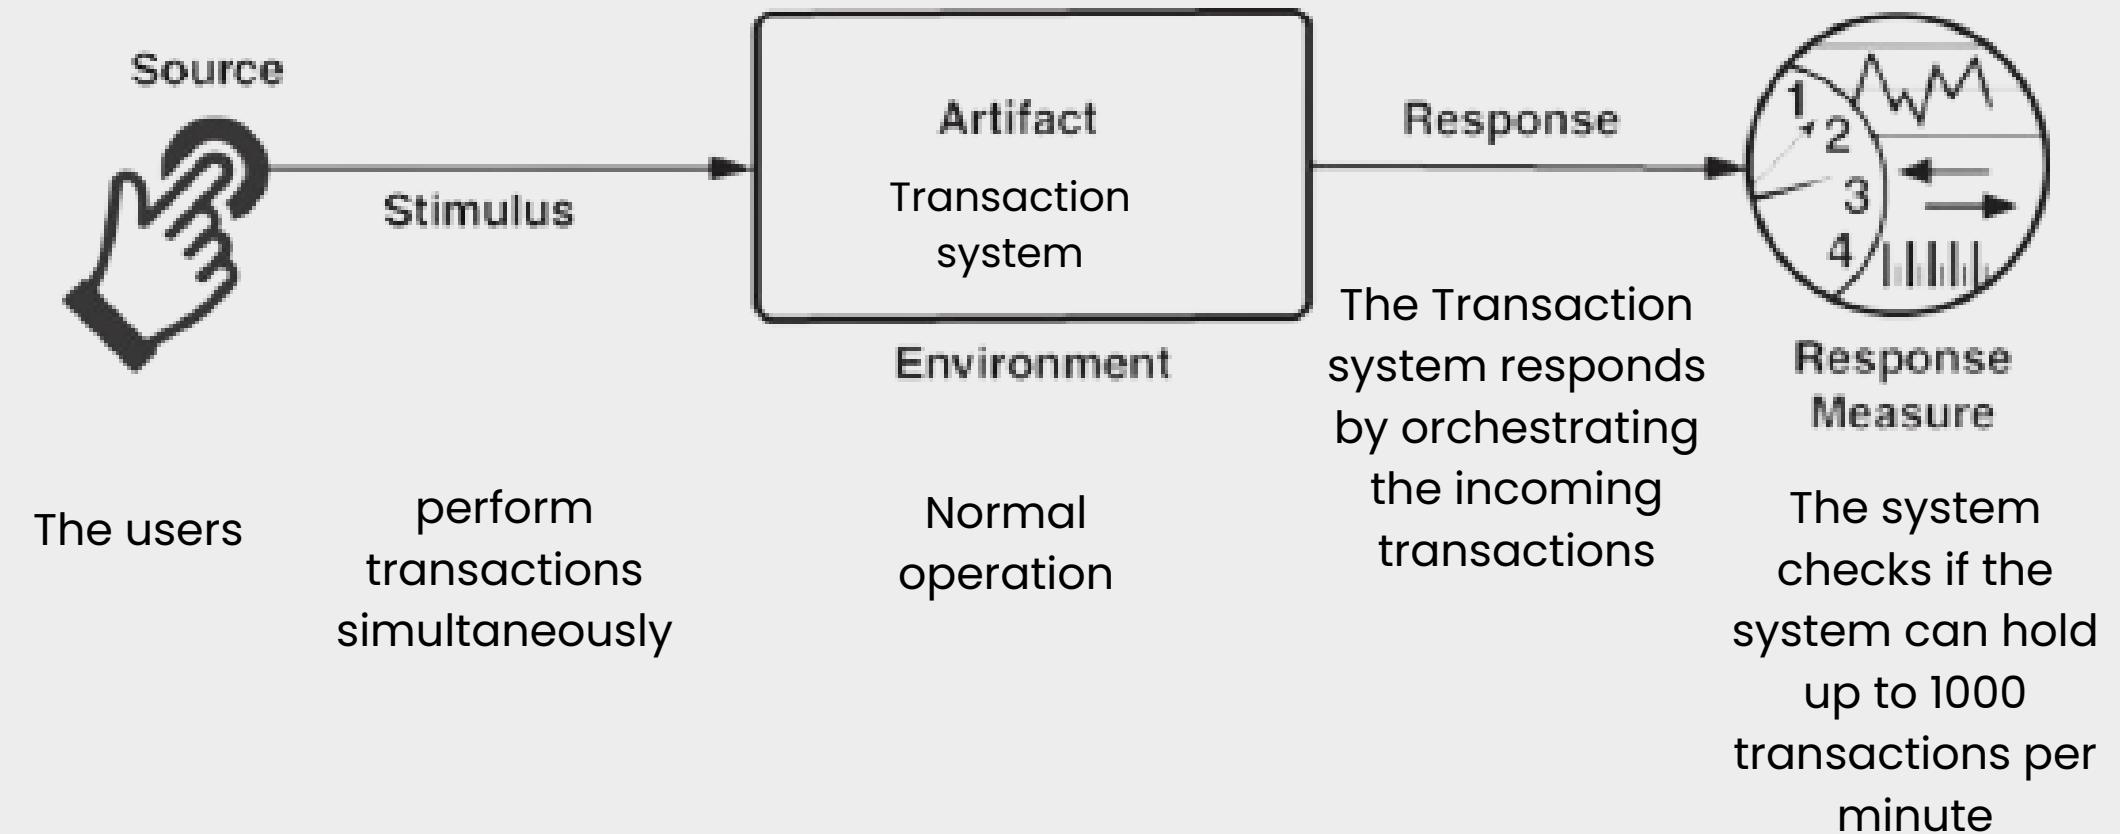


As we are building a social media platform that needed to store and process large amounts of user-generated content, such as photos and videos, a **NoSQL** database like **MongoDB** or **Cassandra** would be a better fit because it can handle large amounts of unstructured data and scale easily to meet the needs of a growing user base.

# Database



- Number of transactions (Performance)



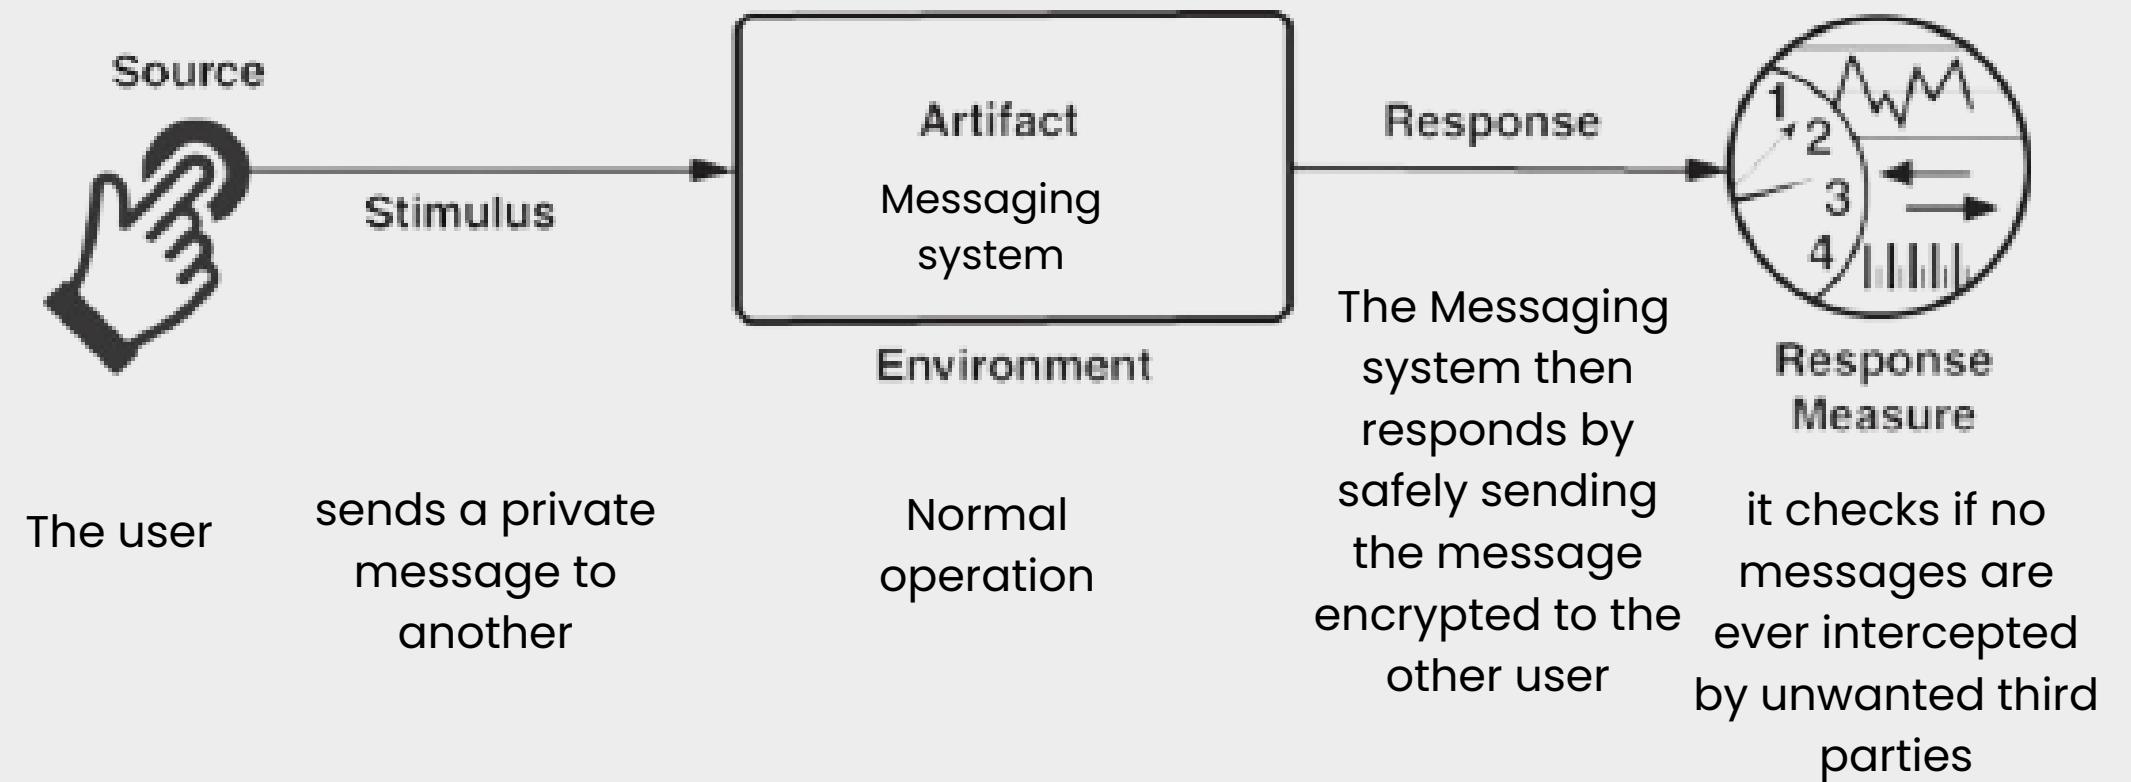
**PostgreSQL** and **MongoDB** are strong contenders for handling high volumes of concurrent transactions, with **PostgreSQL** being a good choice for traditional relational workloads and **MongoDB** being a good choice for handling complex data structures and distributed transactions.

In terms of scalability, **MongoDB** is often considered better because of its horizontal scaling capabilities and sharding feature, which allows for distributing data across multiple machines.

# Database



- Secure messaging (Security)



**PostgreSQL, MongoDB, and Oracle** offer more comprehensive built-in security features, including encryption capabilities. These databases provide better options for ensuring the confidentiality and integrity of the messages.

# Database



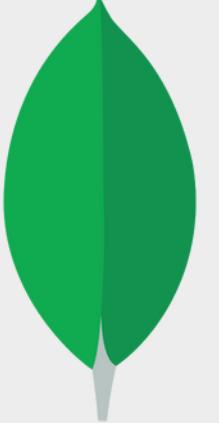
- Low Funds
  - As a Client that is starting a business, I want to save the most I can seeing that I do not have a lot of funds

Considering that we will be the developers of the platform, as mentioned previously in the framework section, in order to save funds and attain a faster development, the most suitable database would be one with which we are familiar, such as **MongoDB** or **PostgreSQL**.

# Database

Decision



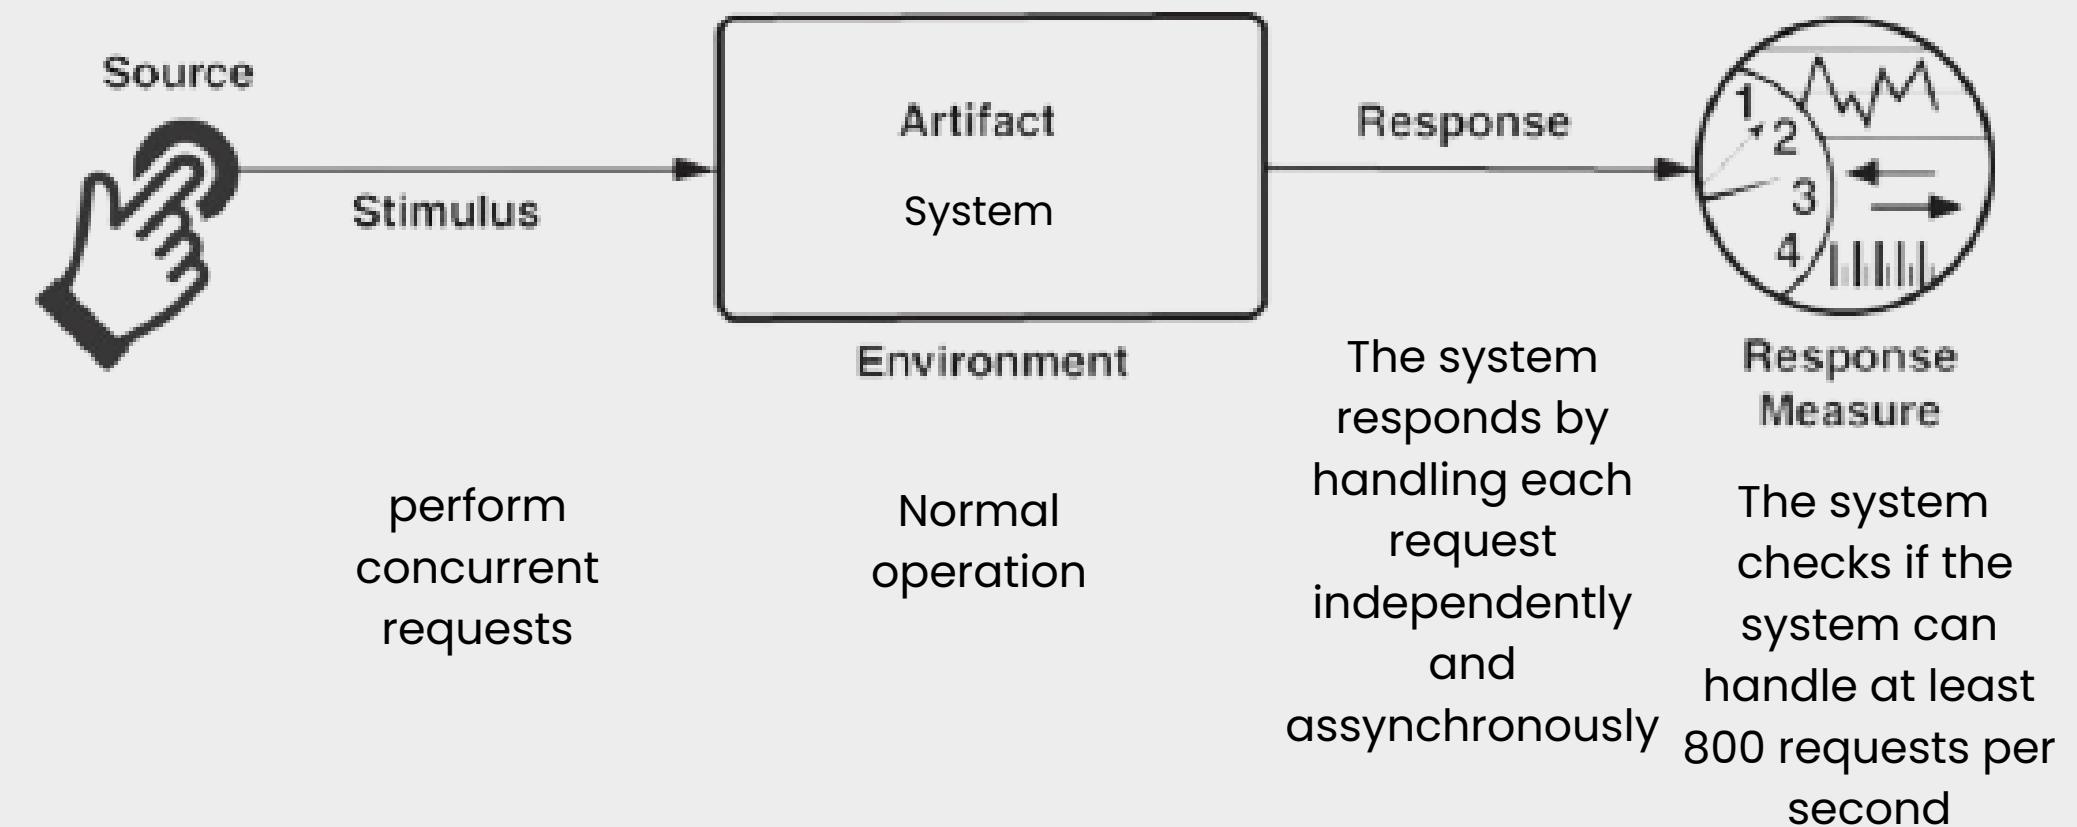
 mongoDB®

The MongoDB logo, featuring a green leaf icon followed by the word "mongoDB" in a lowercase sans-serif font, with a registered trademark symbol (®) at the end.

# Architecture Type



- Number of concurrent requests (Performance)

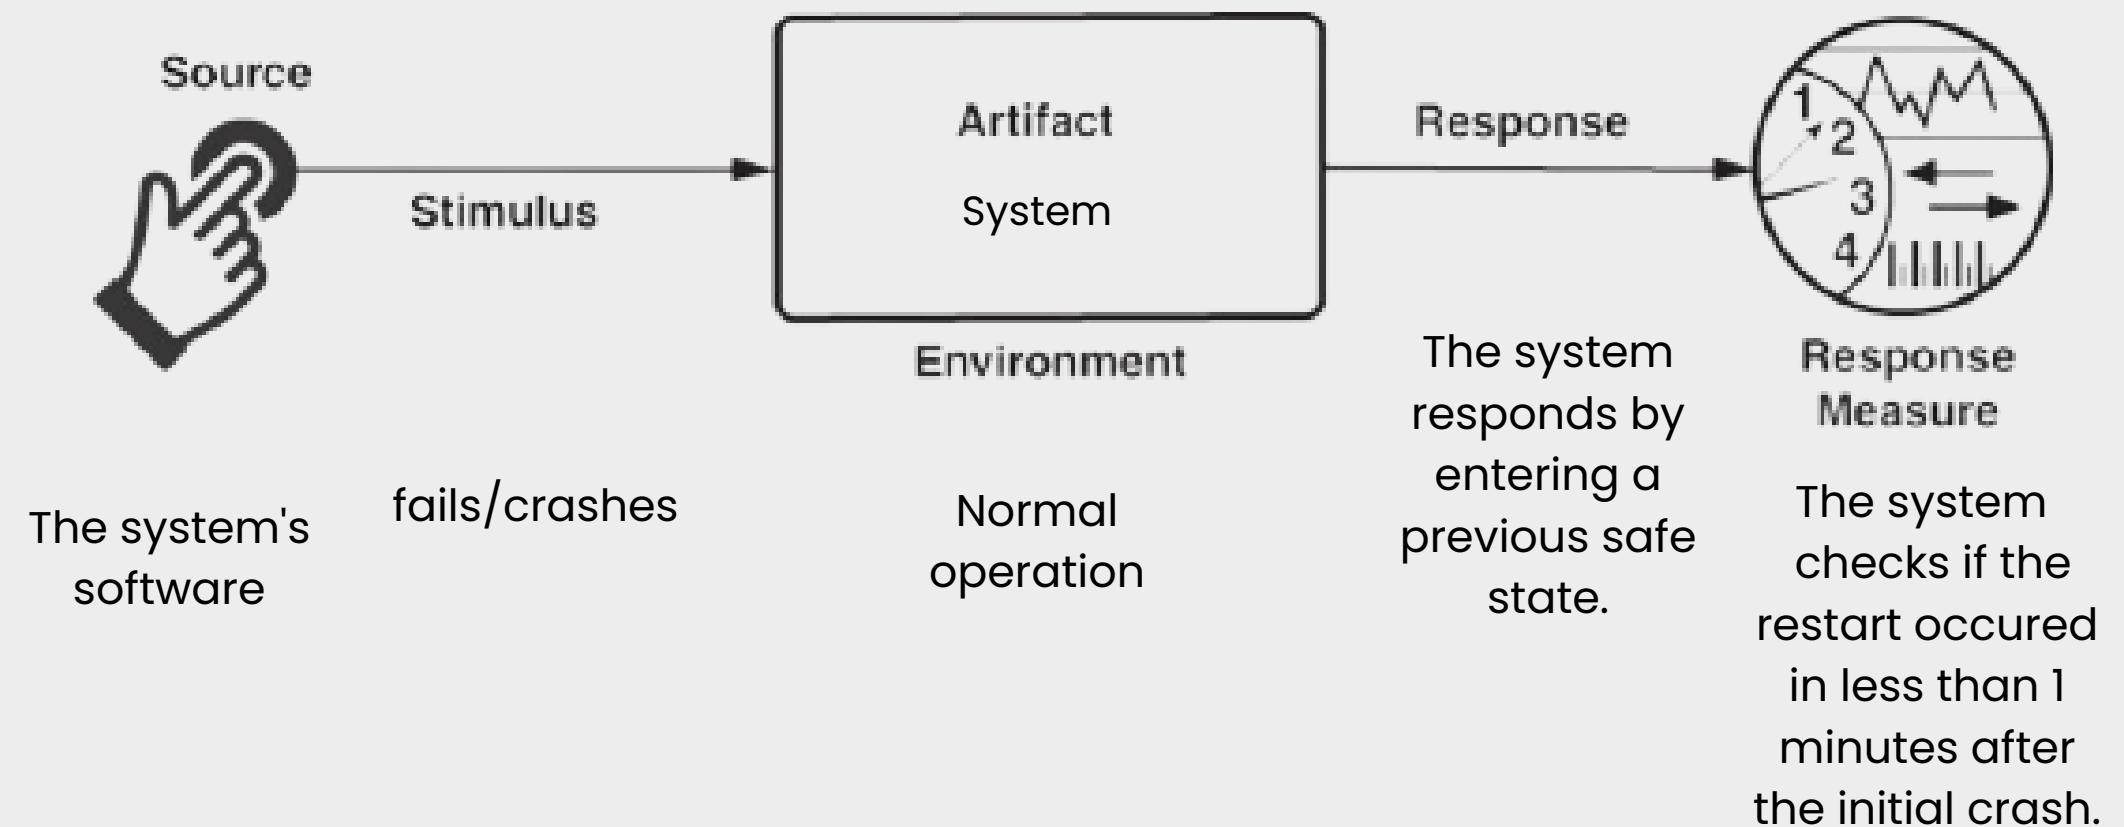


**Event-driven architecture** and **space-based architecture**

# Architecture Type



- Up time (Availability)

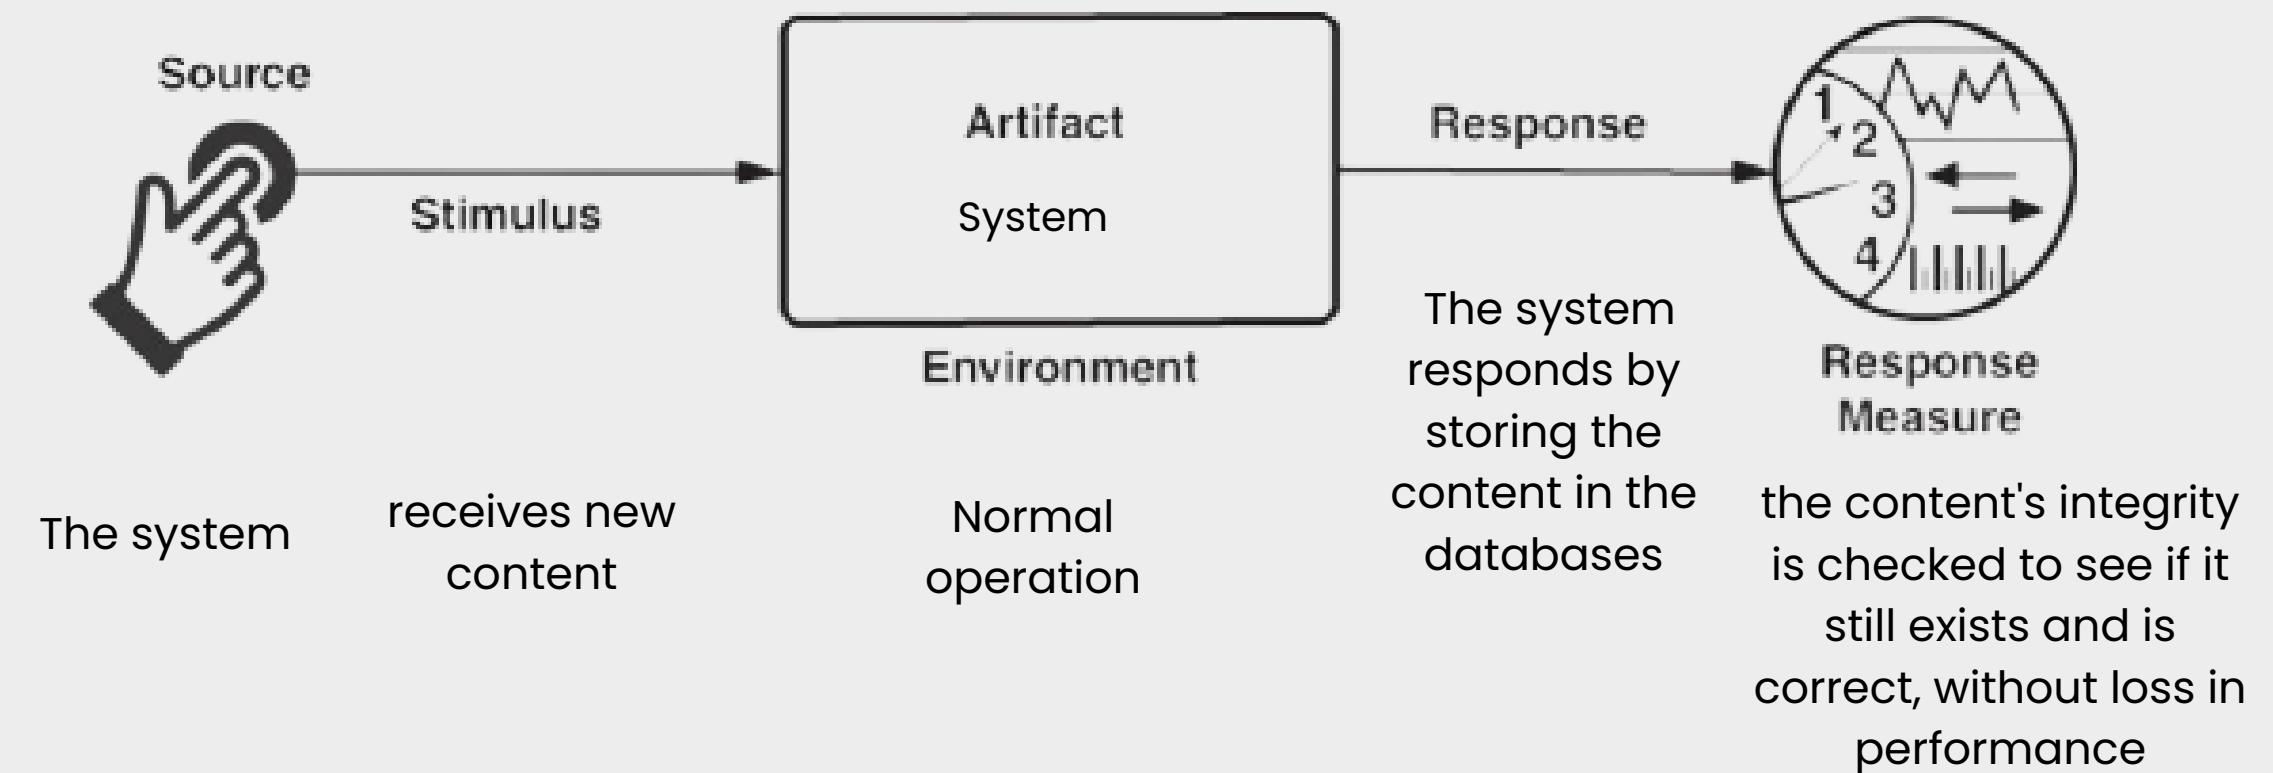


A **microservices architecture** or **service-based architecture** can be beneficial since these architectures can provide fault isolation and allow for components to be restarted independently, reducing the chances of a full system failure.

# Architecture Type



- Sustain an annual content growth of at least 200% in content (Scalability)

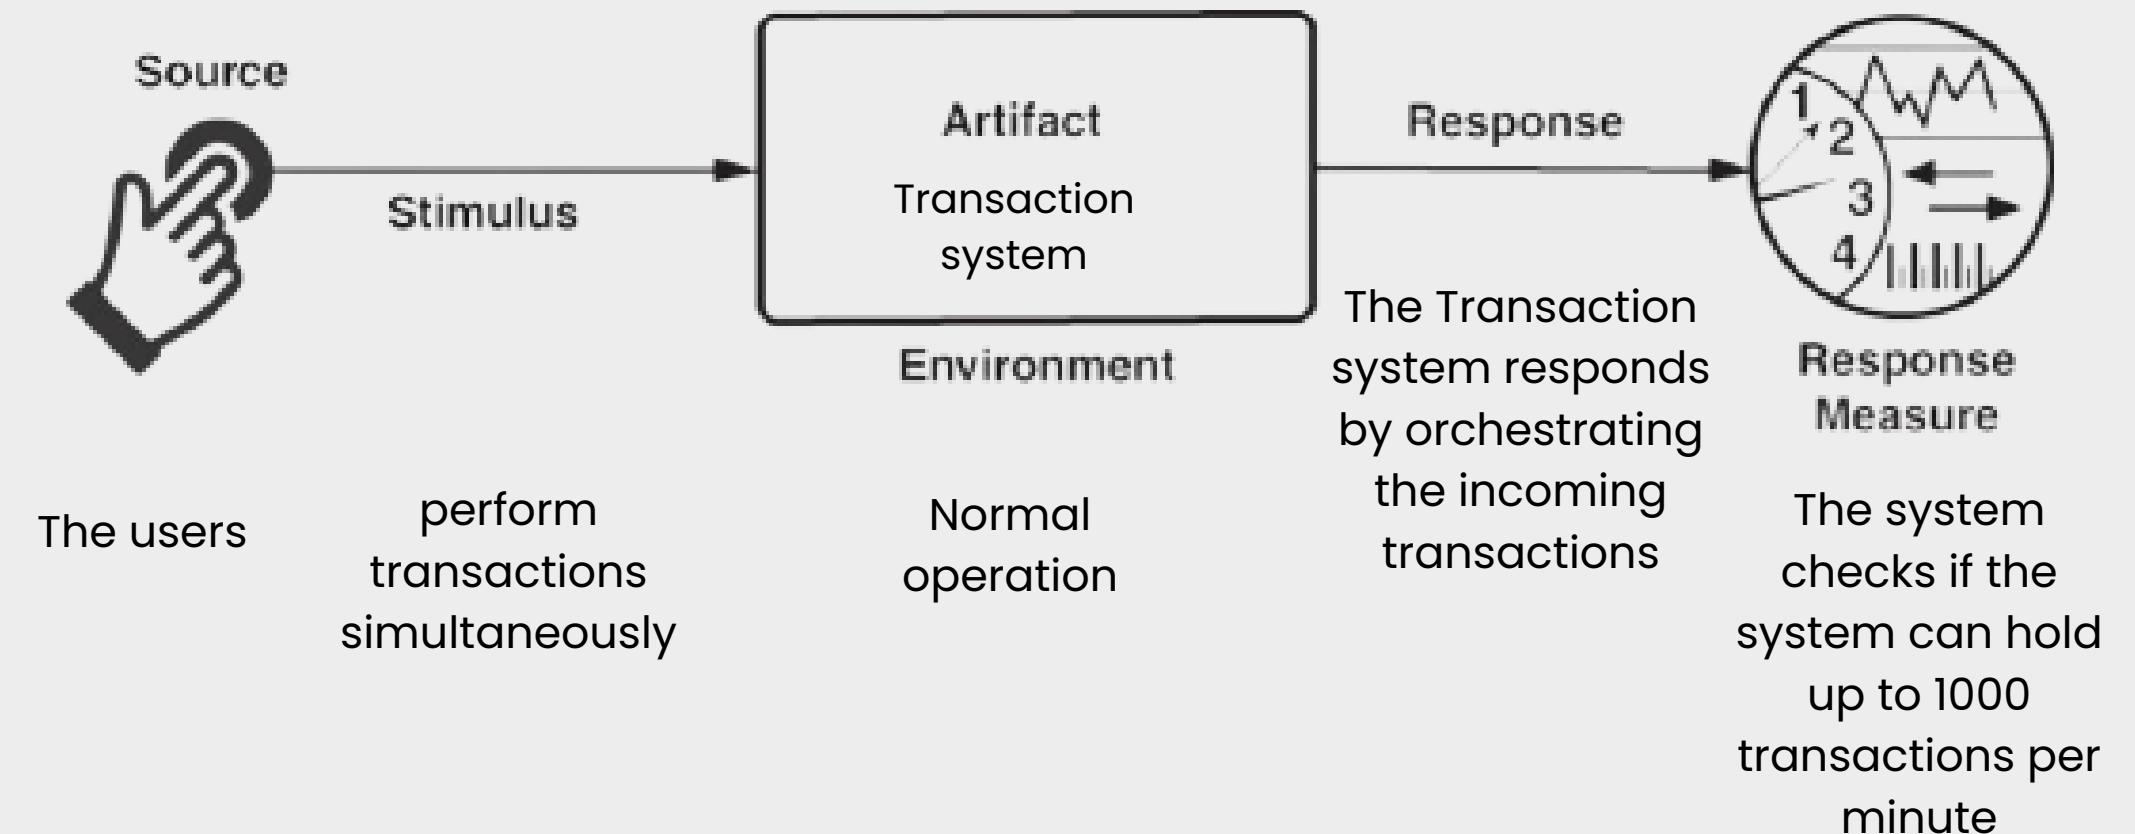


**Event-driven architecture** and **space-based architecture**

# Architecture Type



- Number of transactions (Performance)

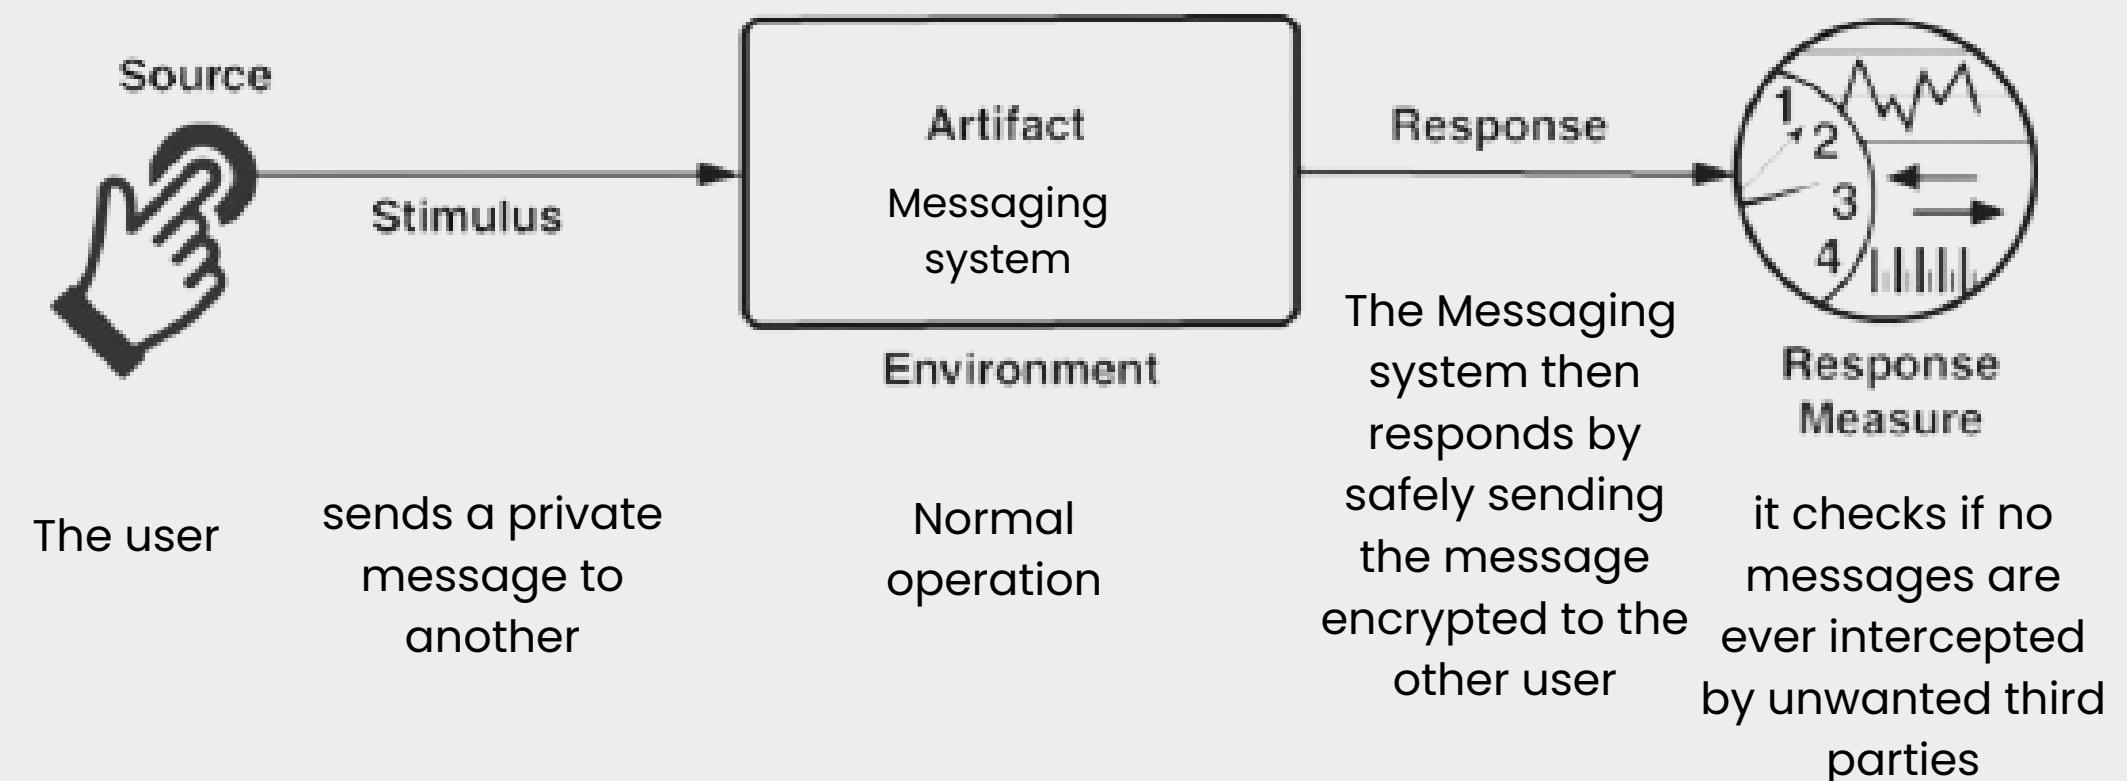


**Event-driven architecture and space-based architecture**

# Architecture Type



- Secure messaging (Security)



**Service-Based Architecture** or a **Microservices Architecture**. In both of these architectures, each service is responsible for performing a specific task, and communication between services is done through APIs, which can be secured using encryption and authentication.

# Architecture Type

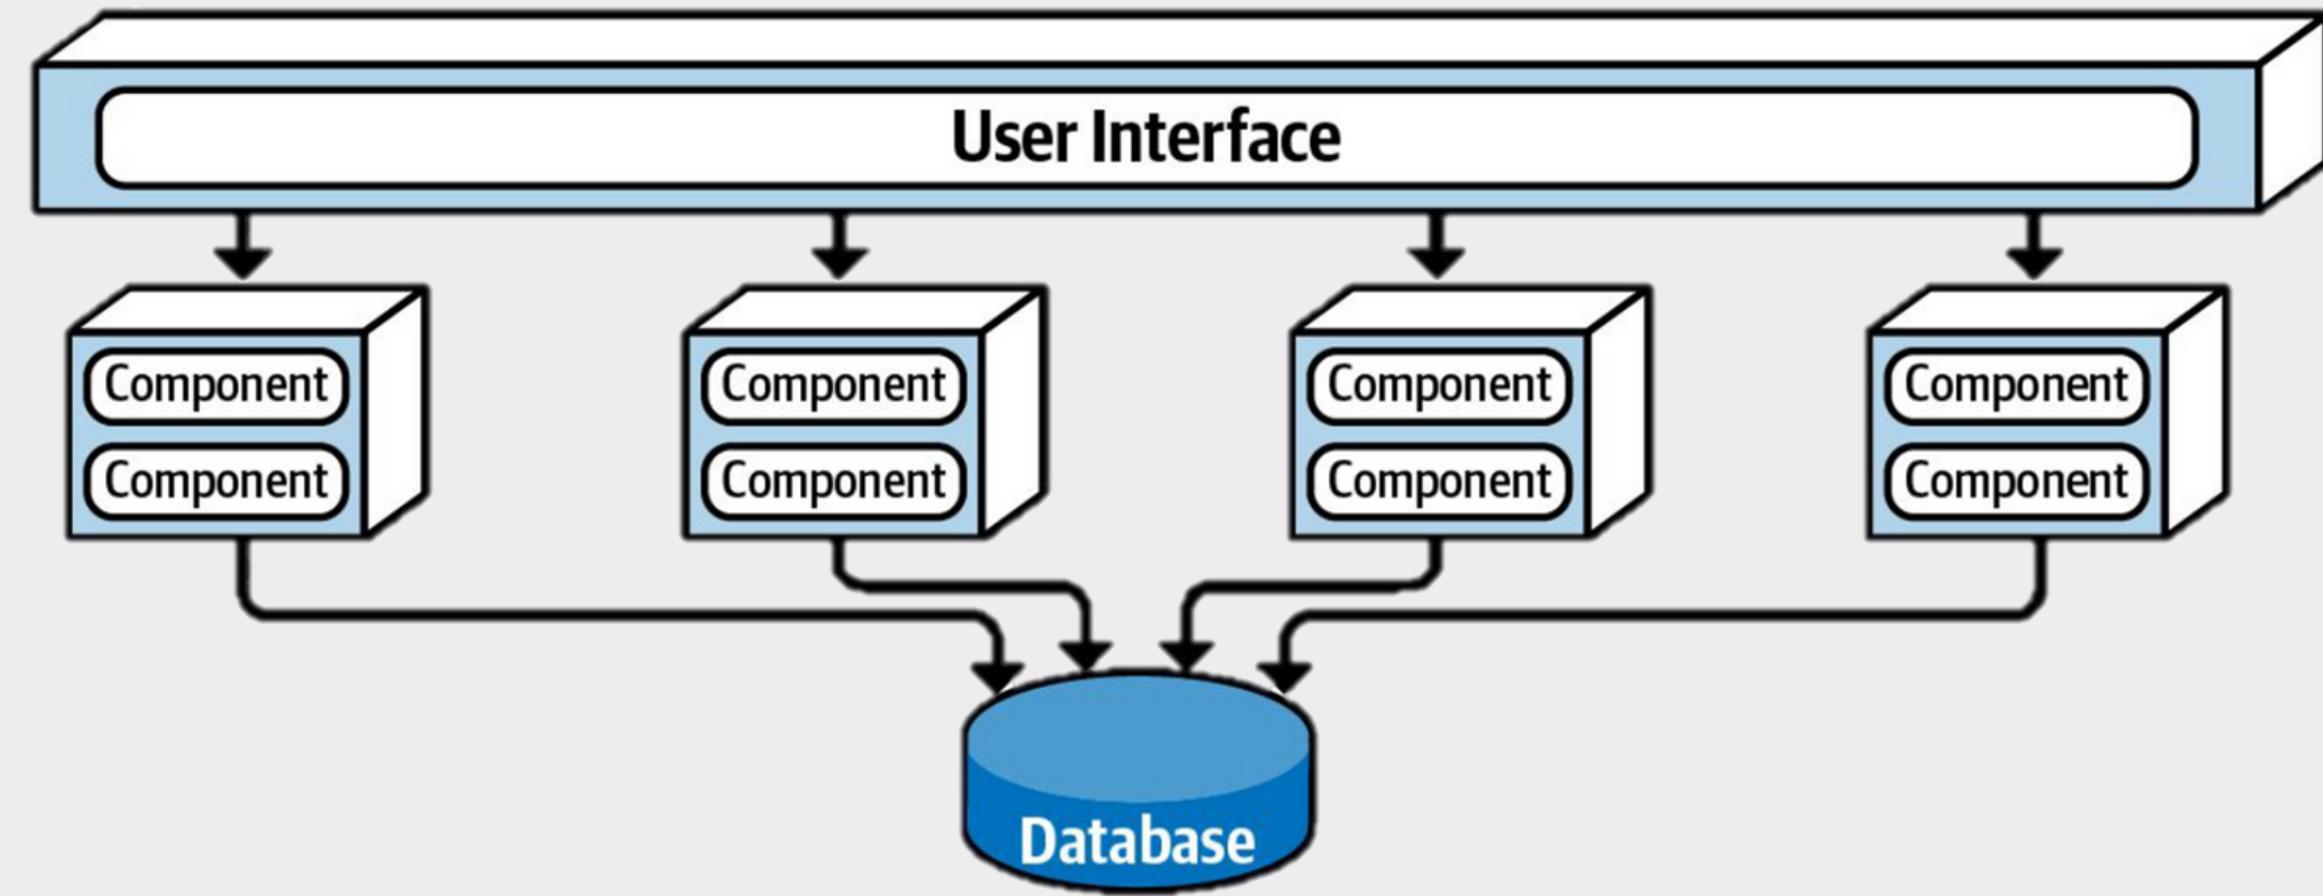


- Low Funds
  - As a Client that is starting a business, I want to save the most I can seeing that I do not have a lot of funds

**Layered architecture, pipelined architecture, microkernel architecture  
and Service-based architecture**

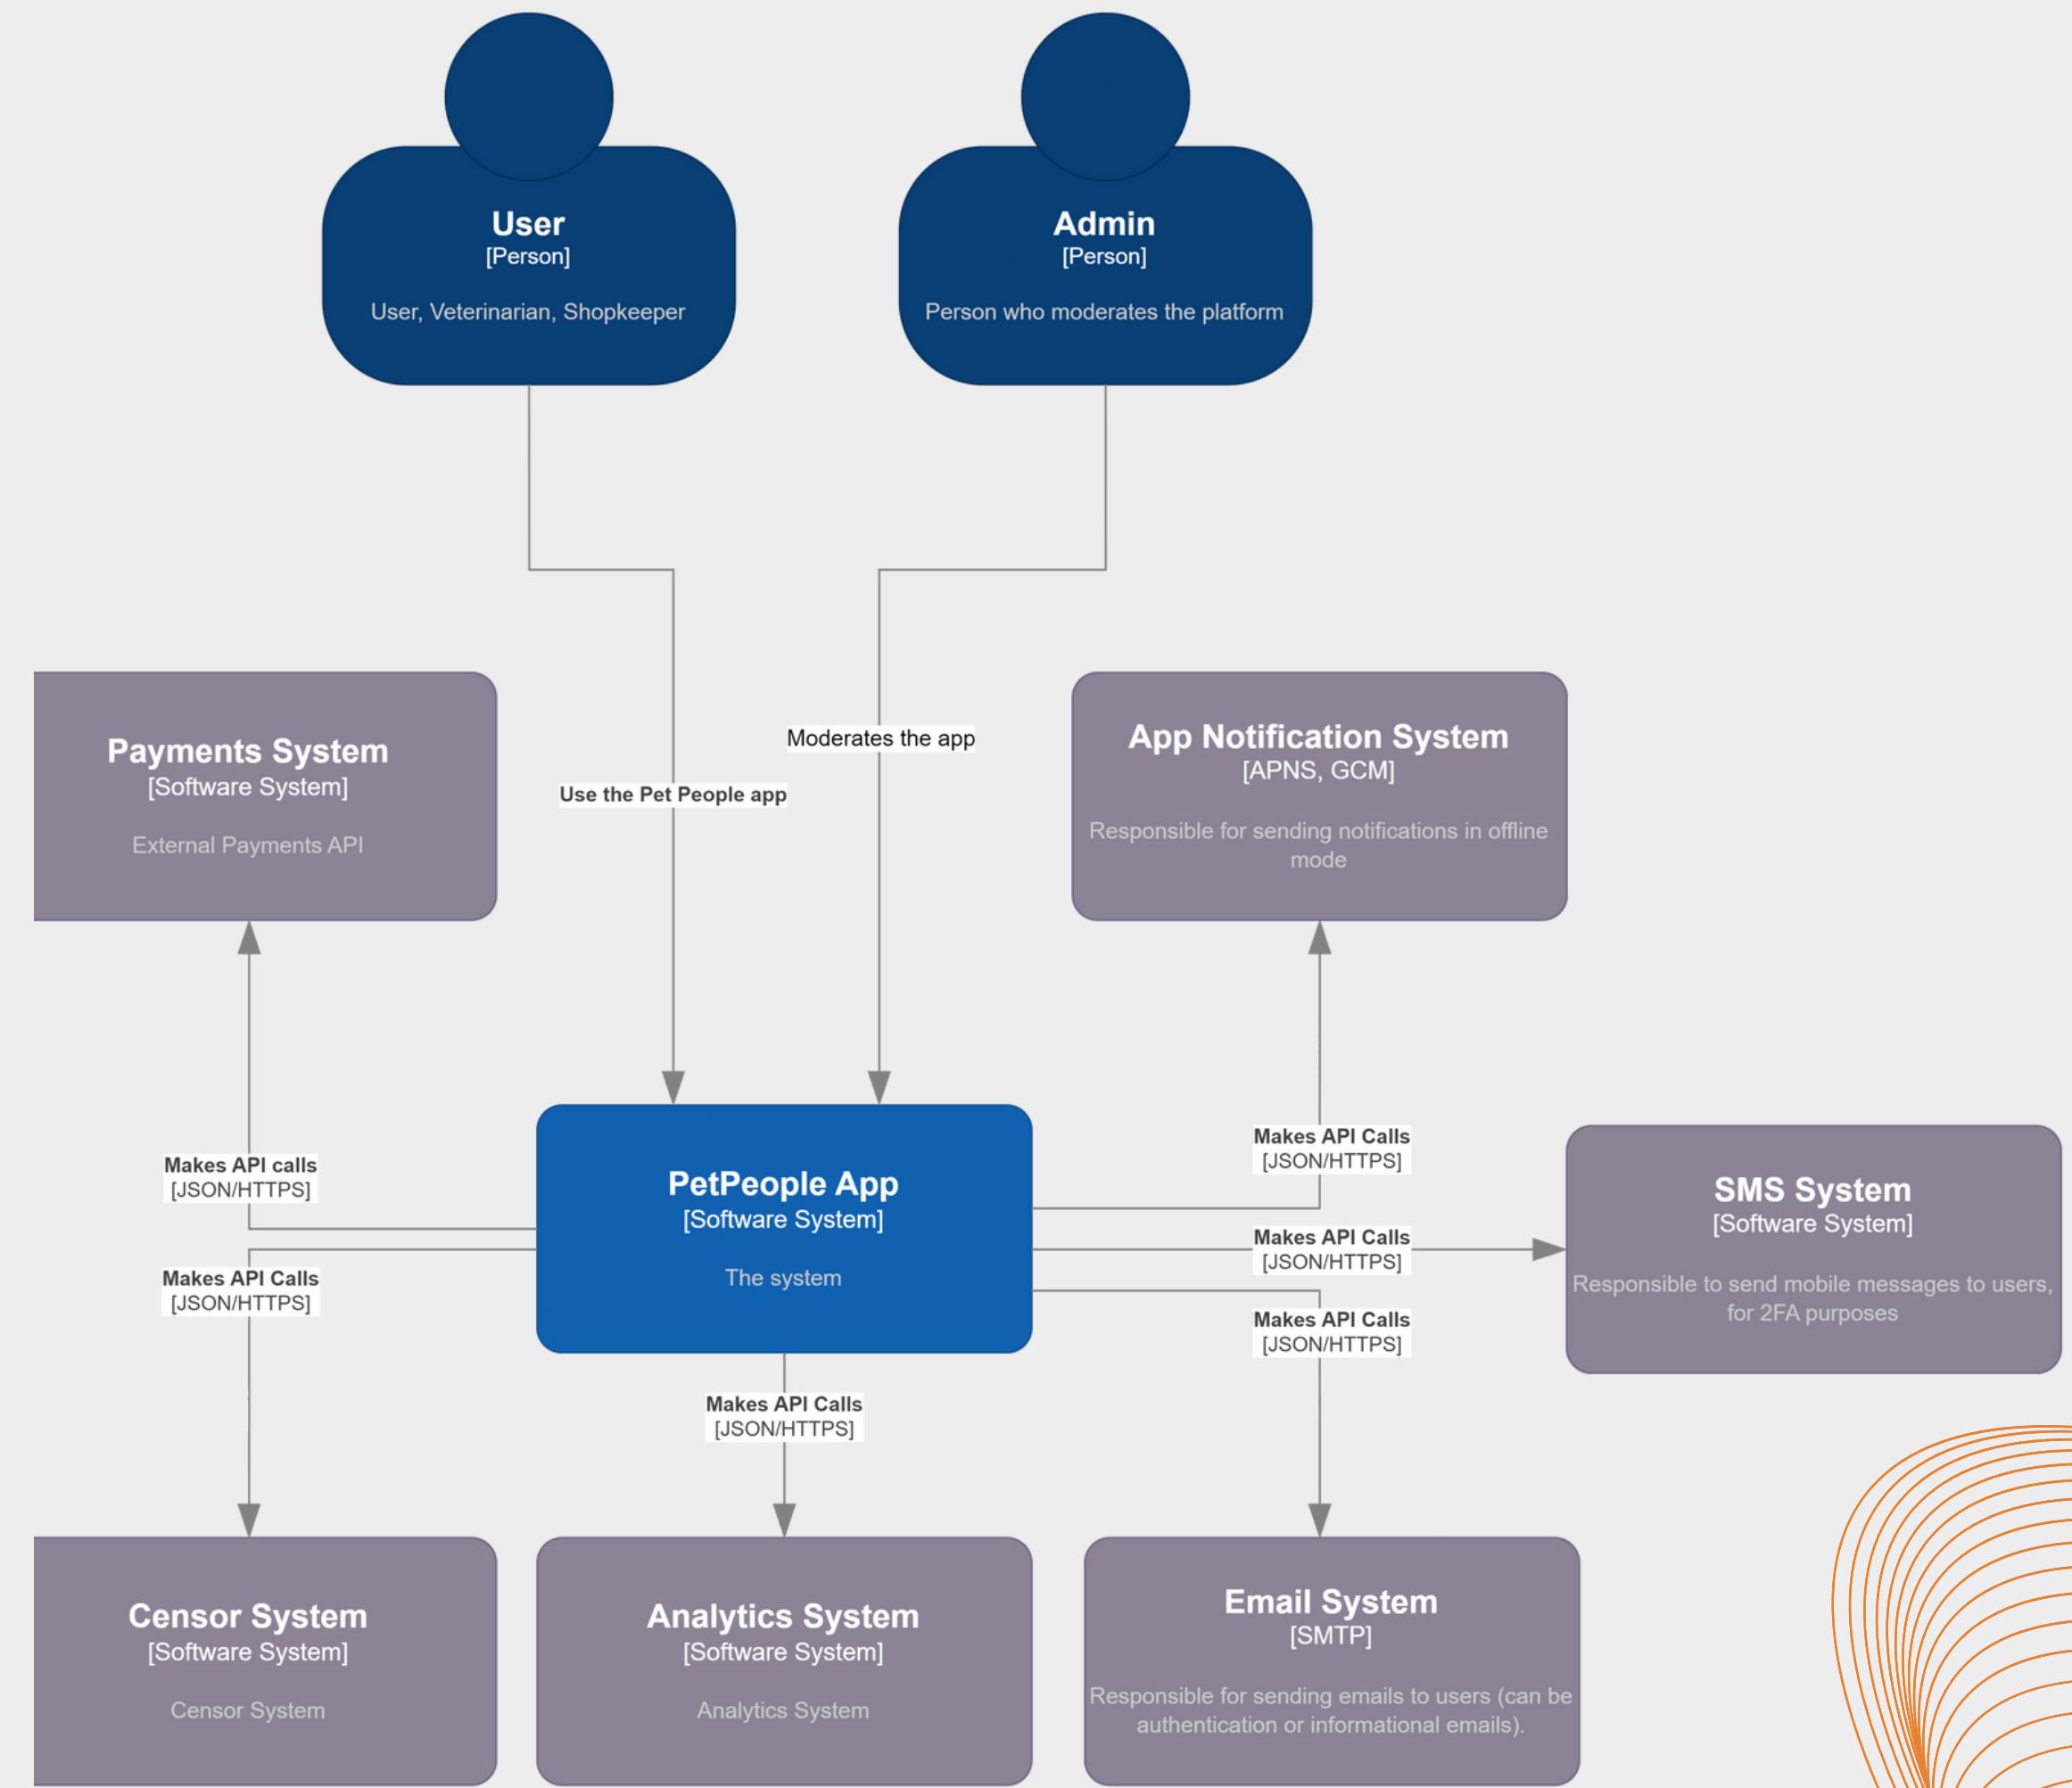
# Architecture Type

## Decision

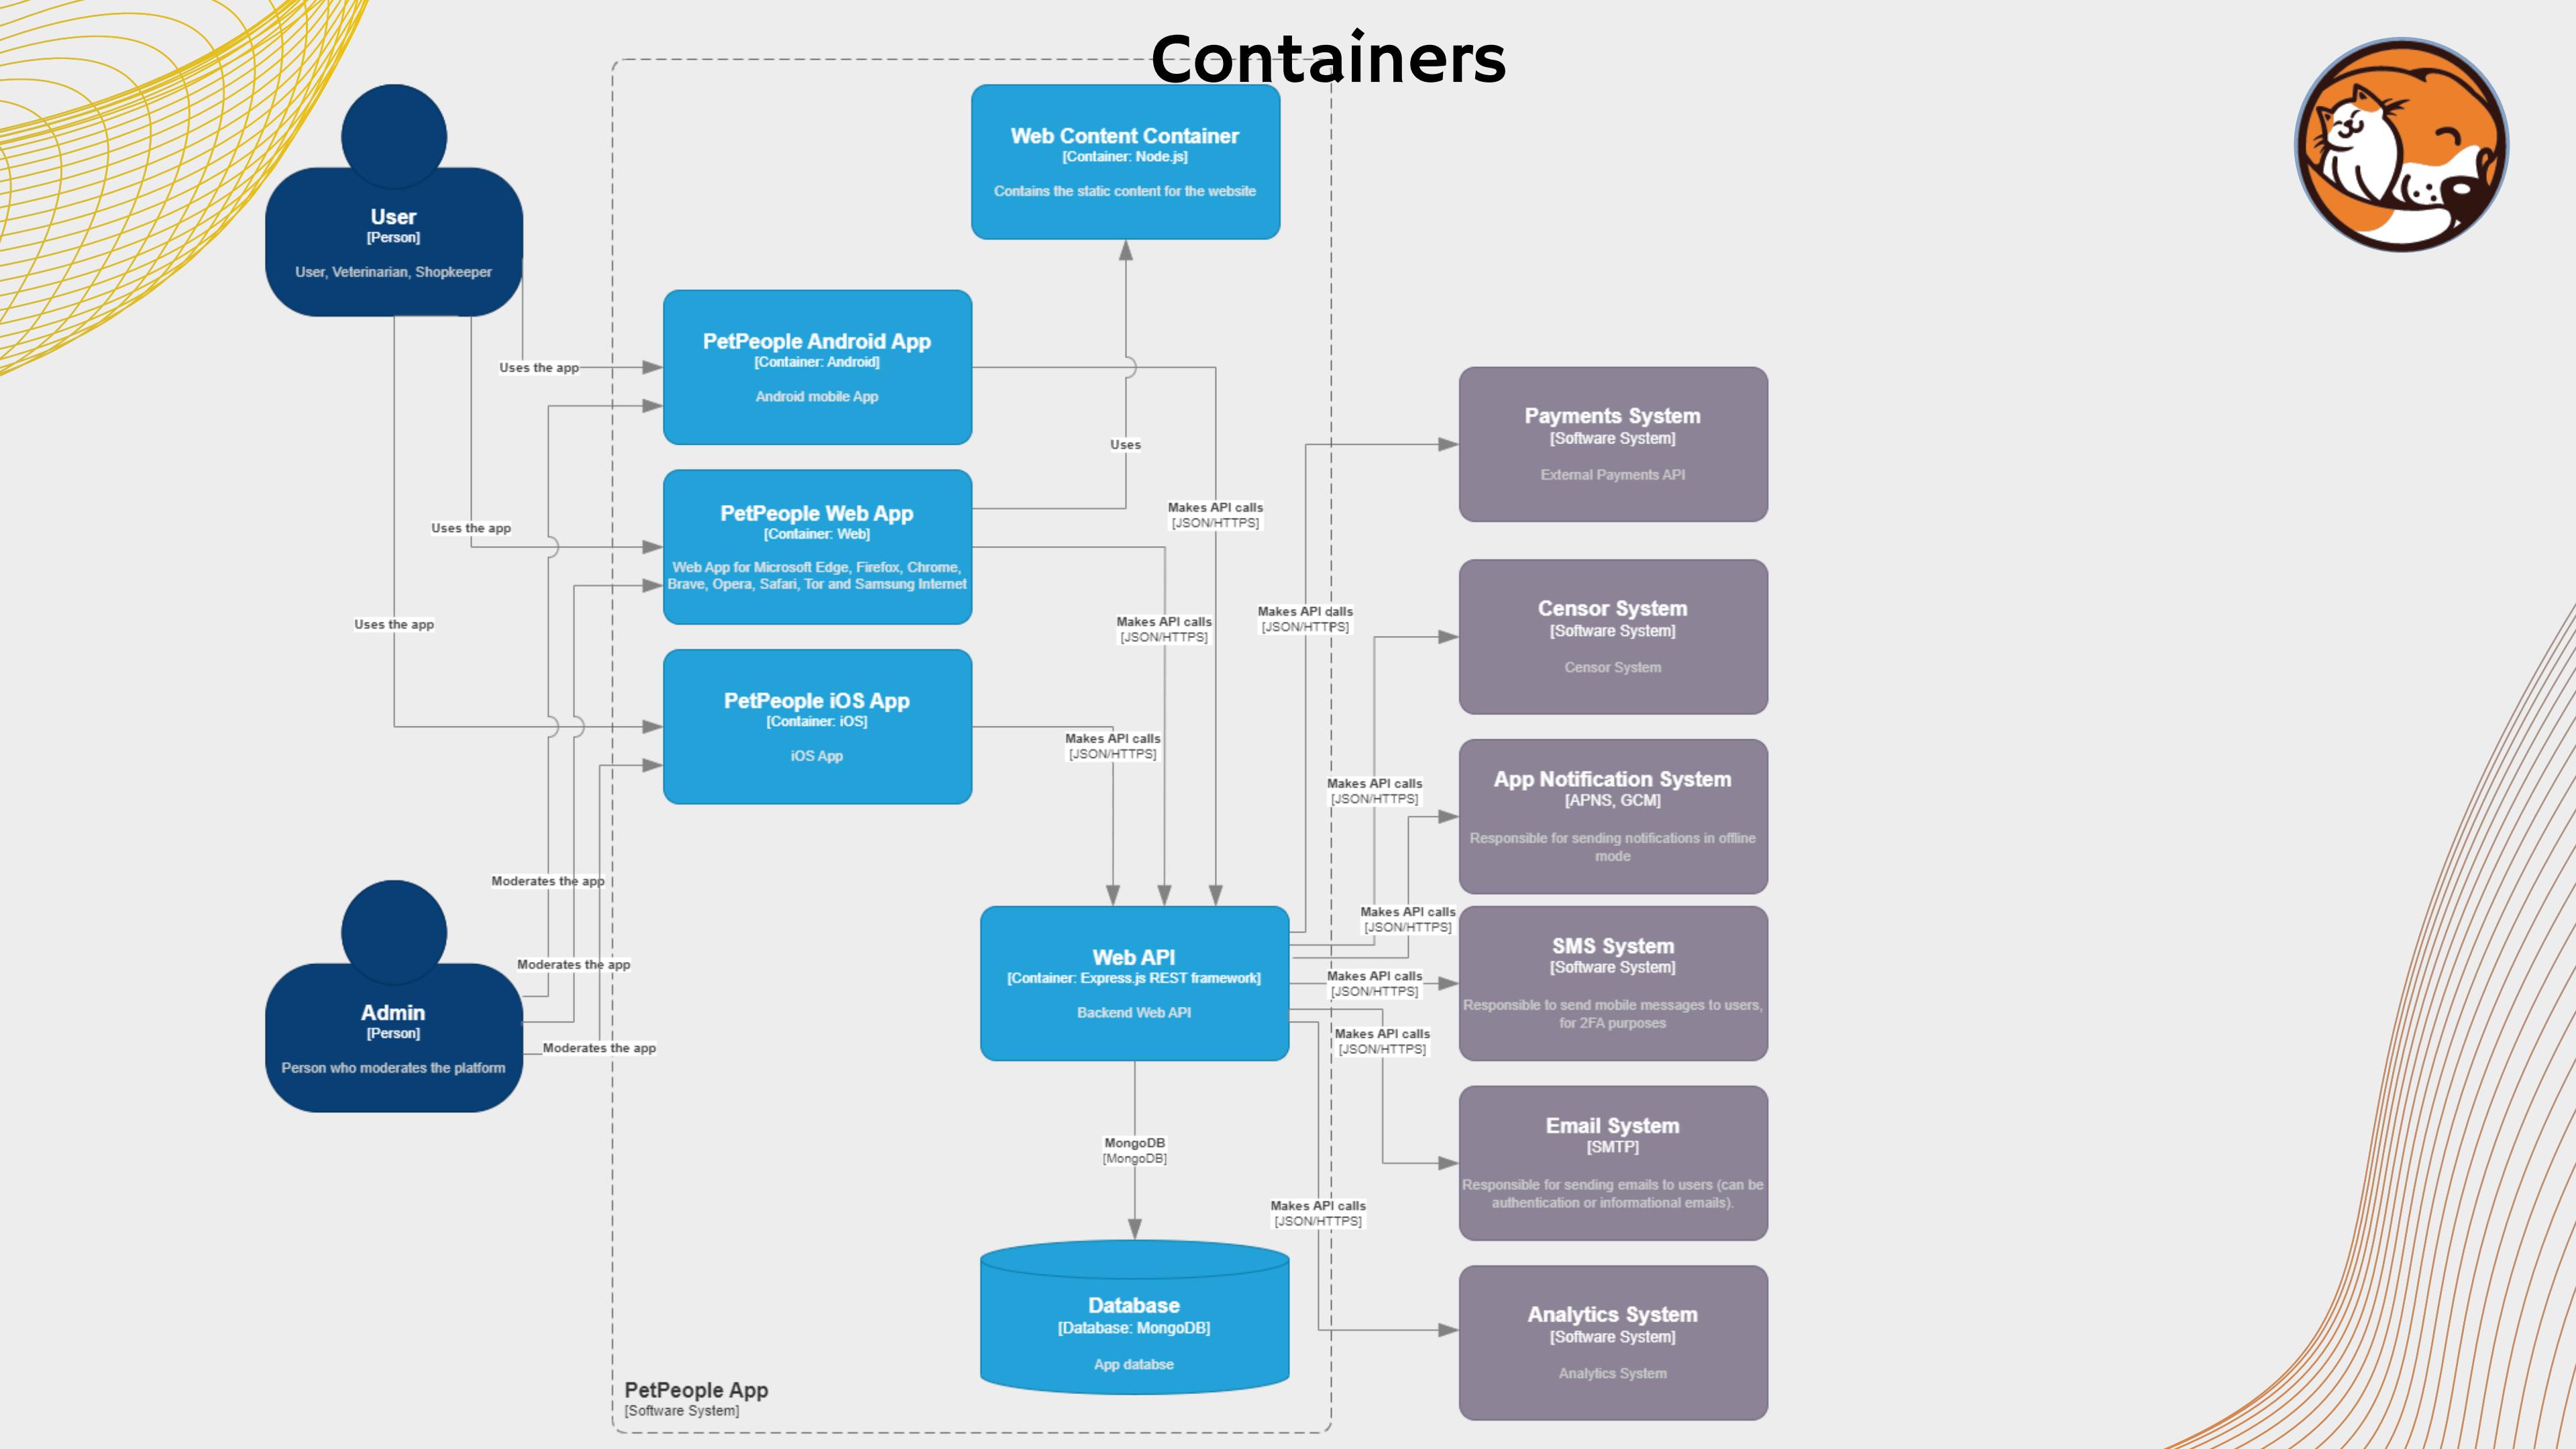


Service-Based architecture

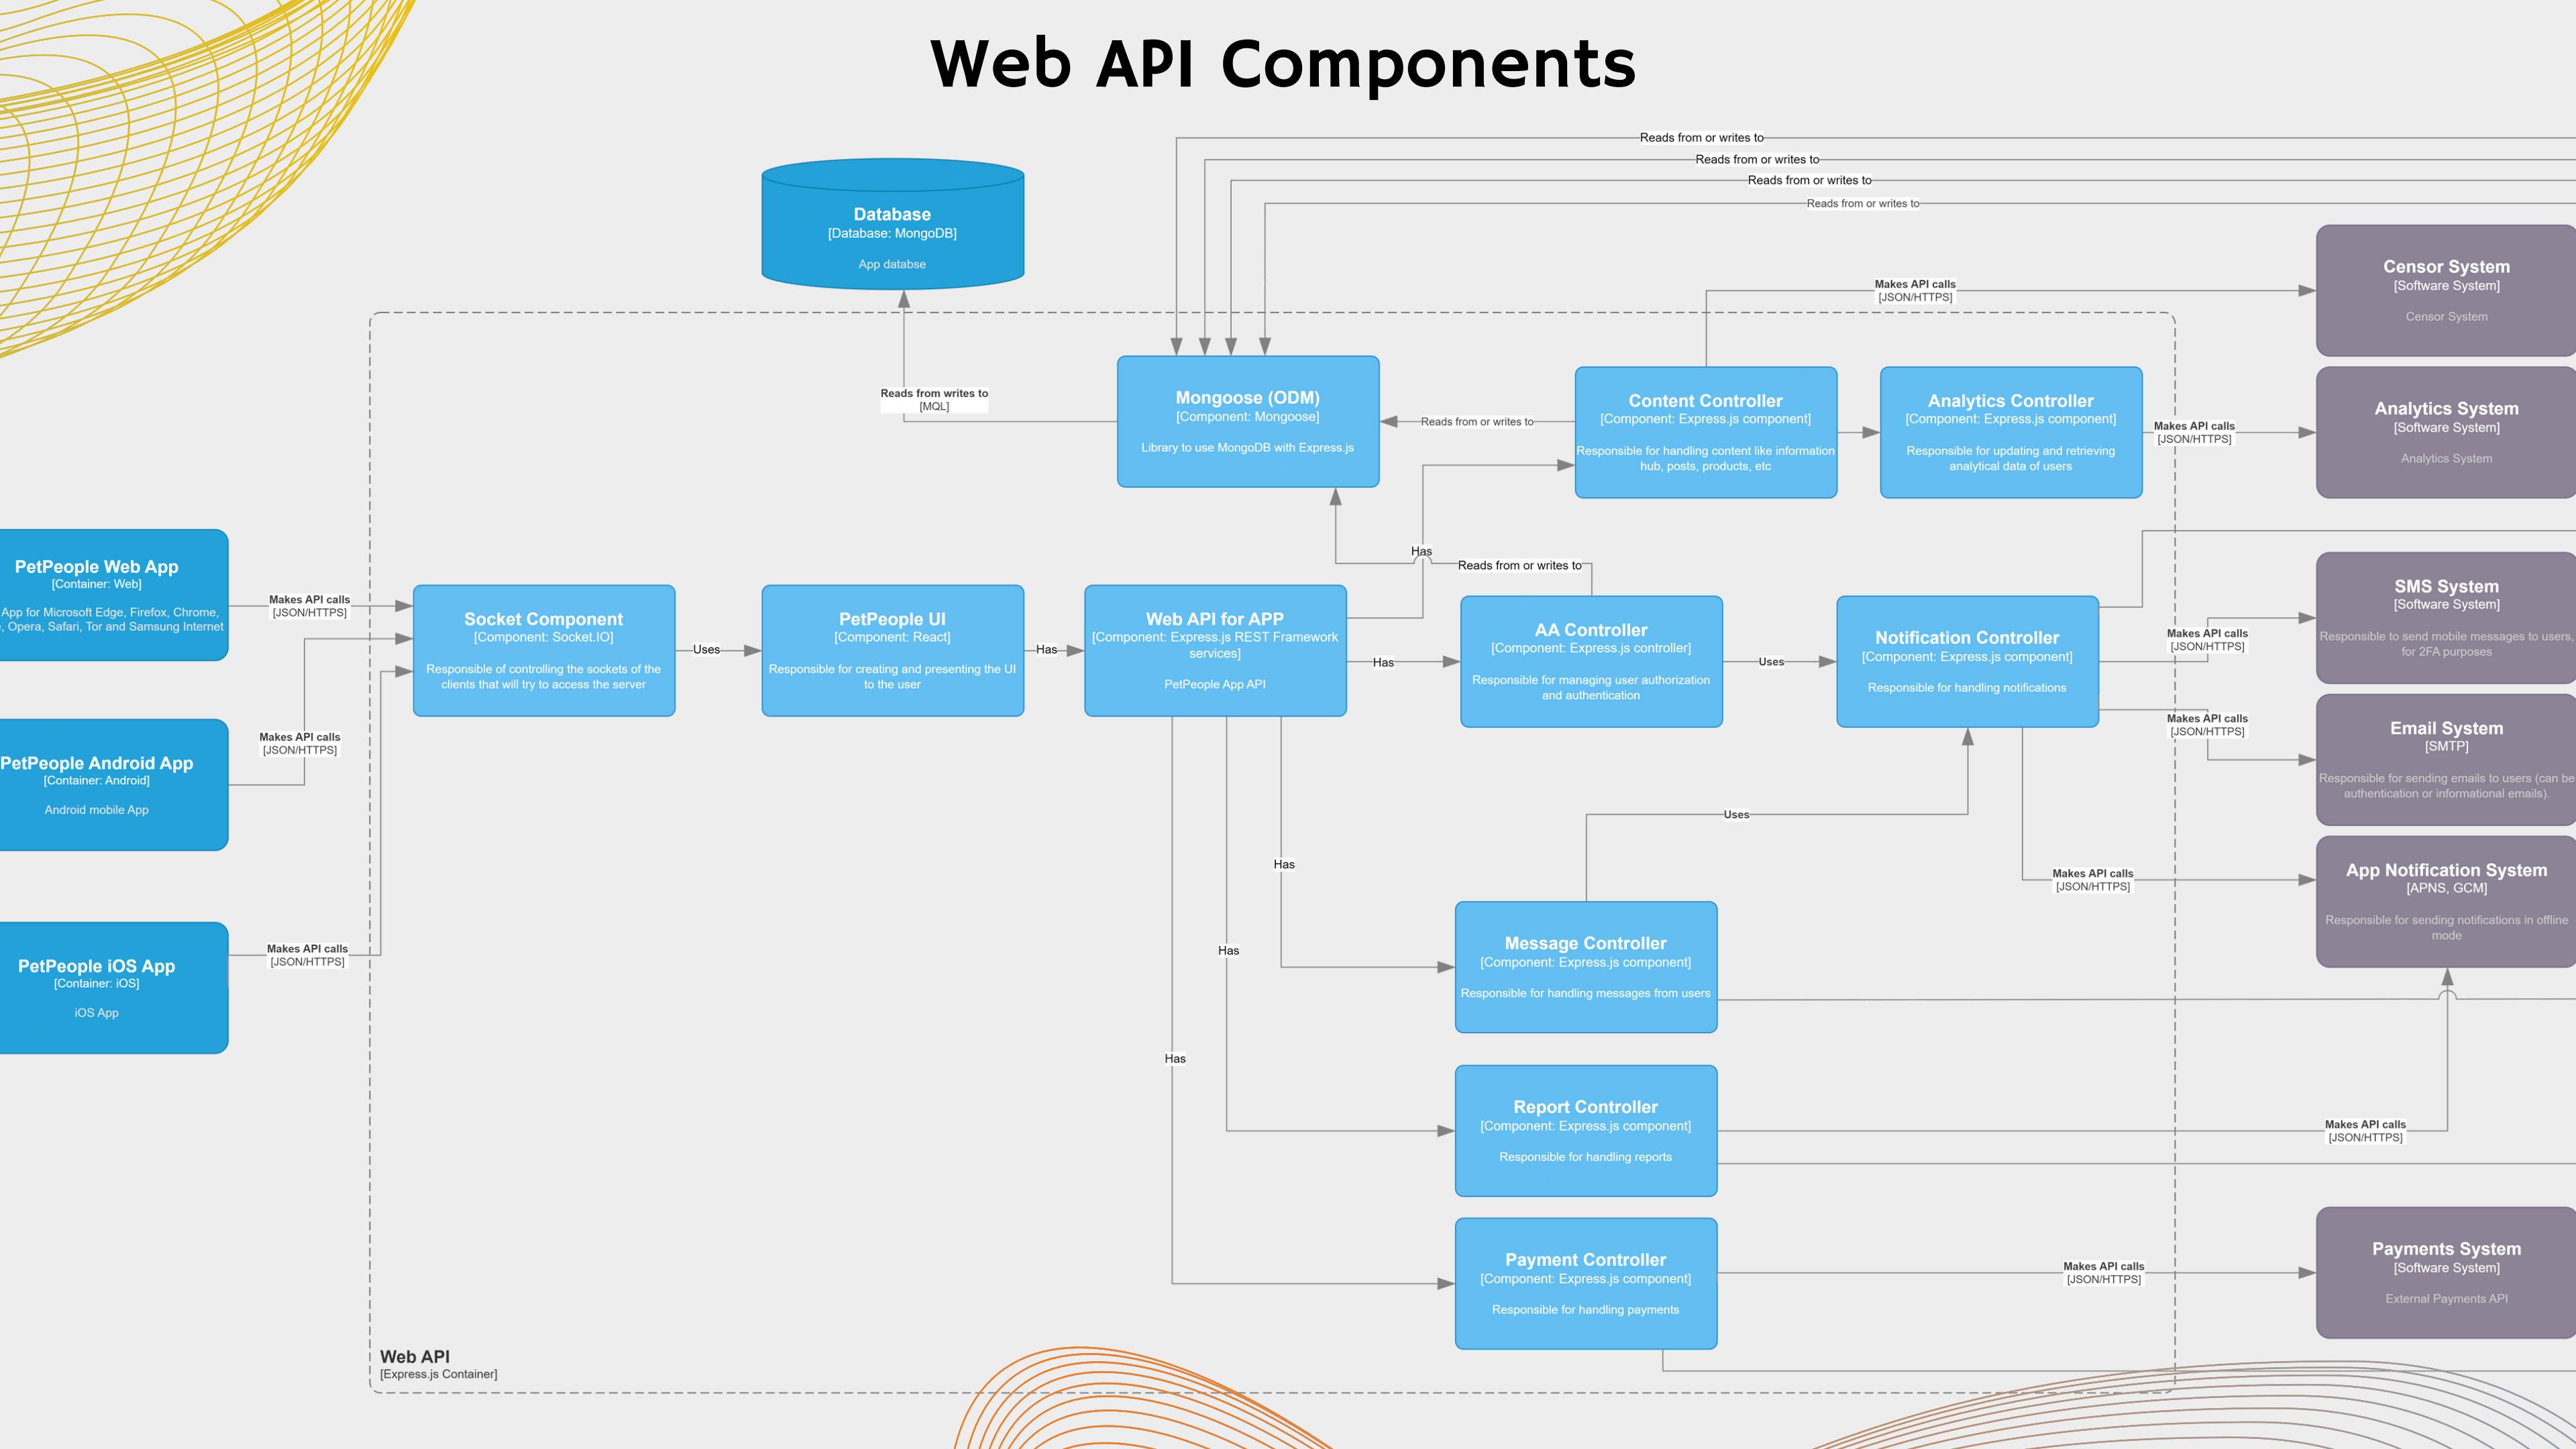
# Context



# Containers



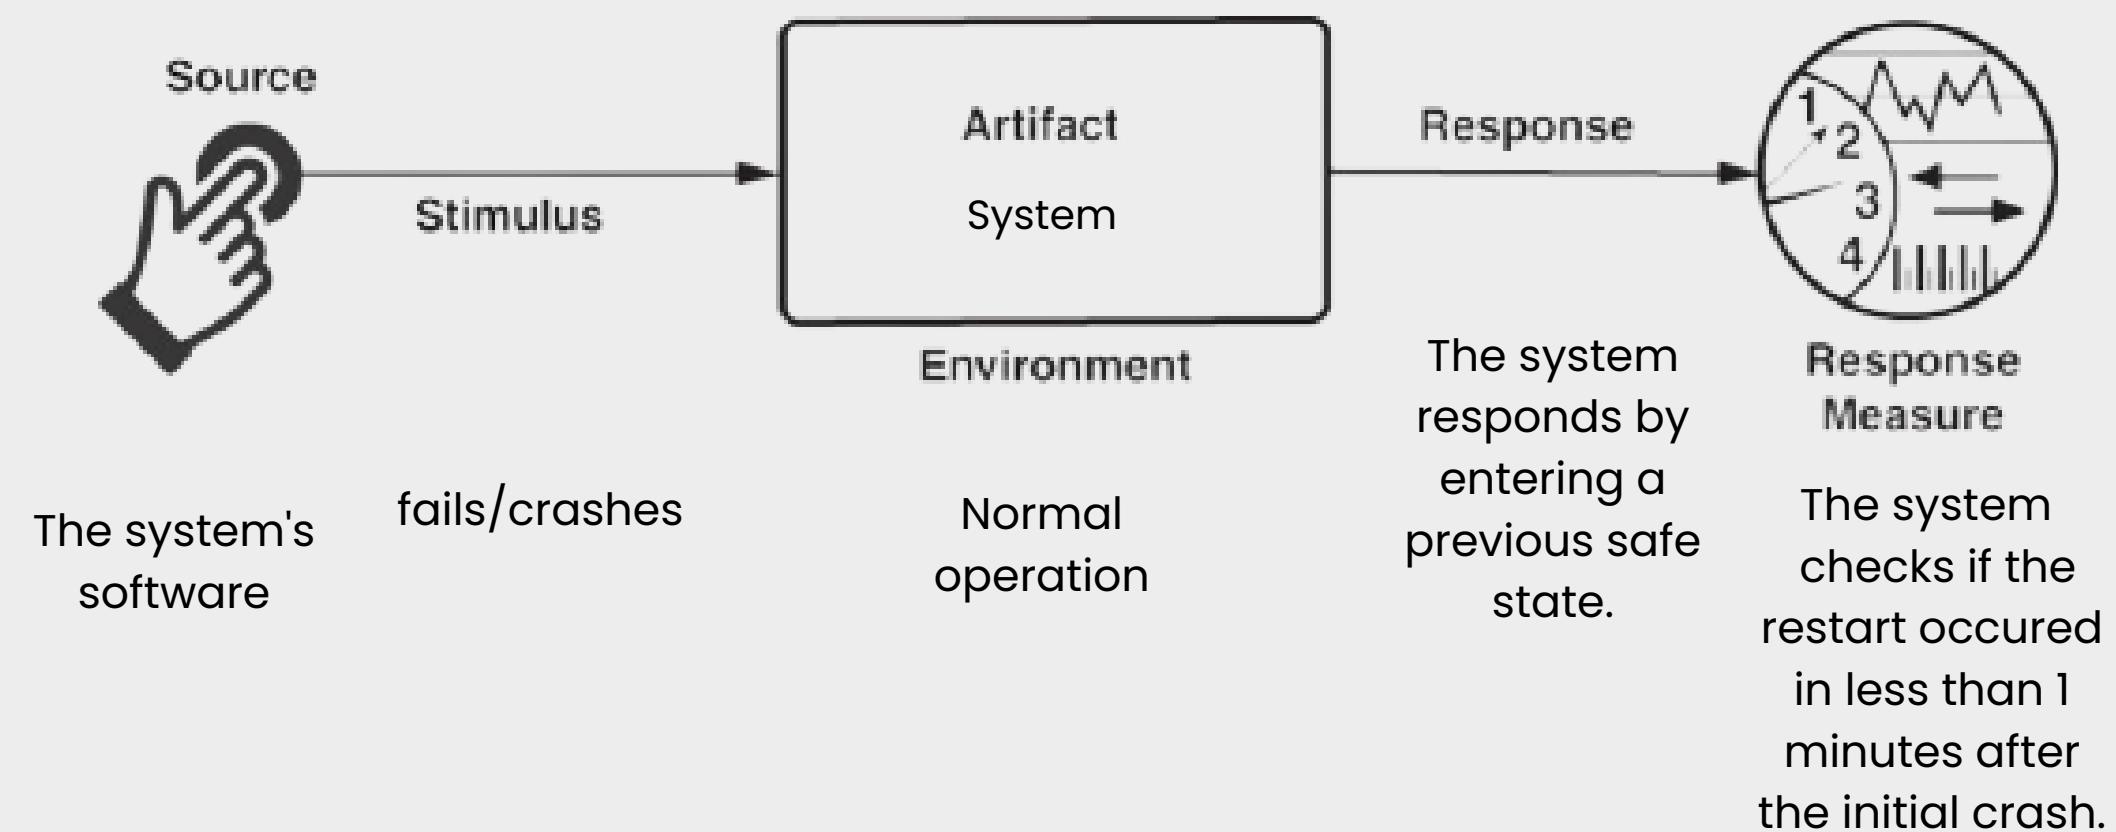
# Web API Components



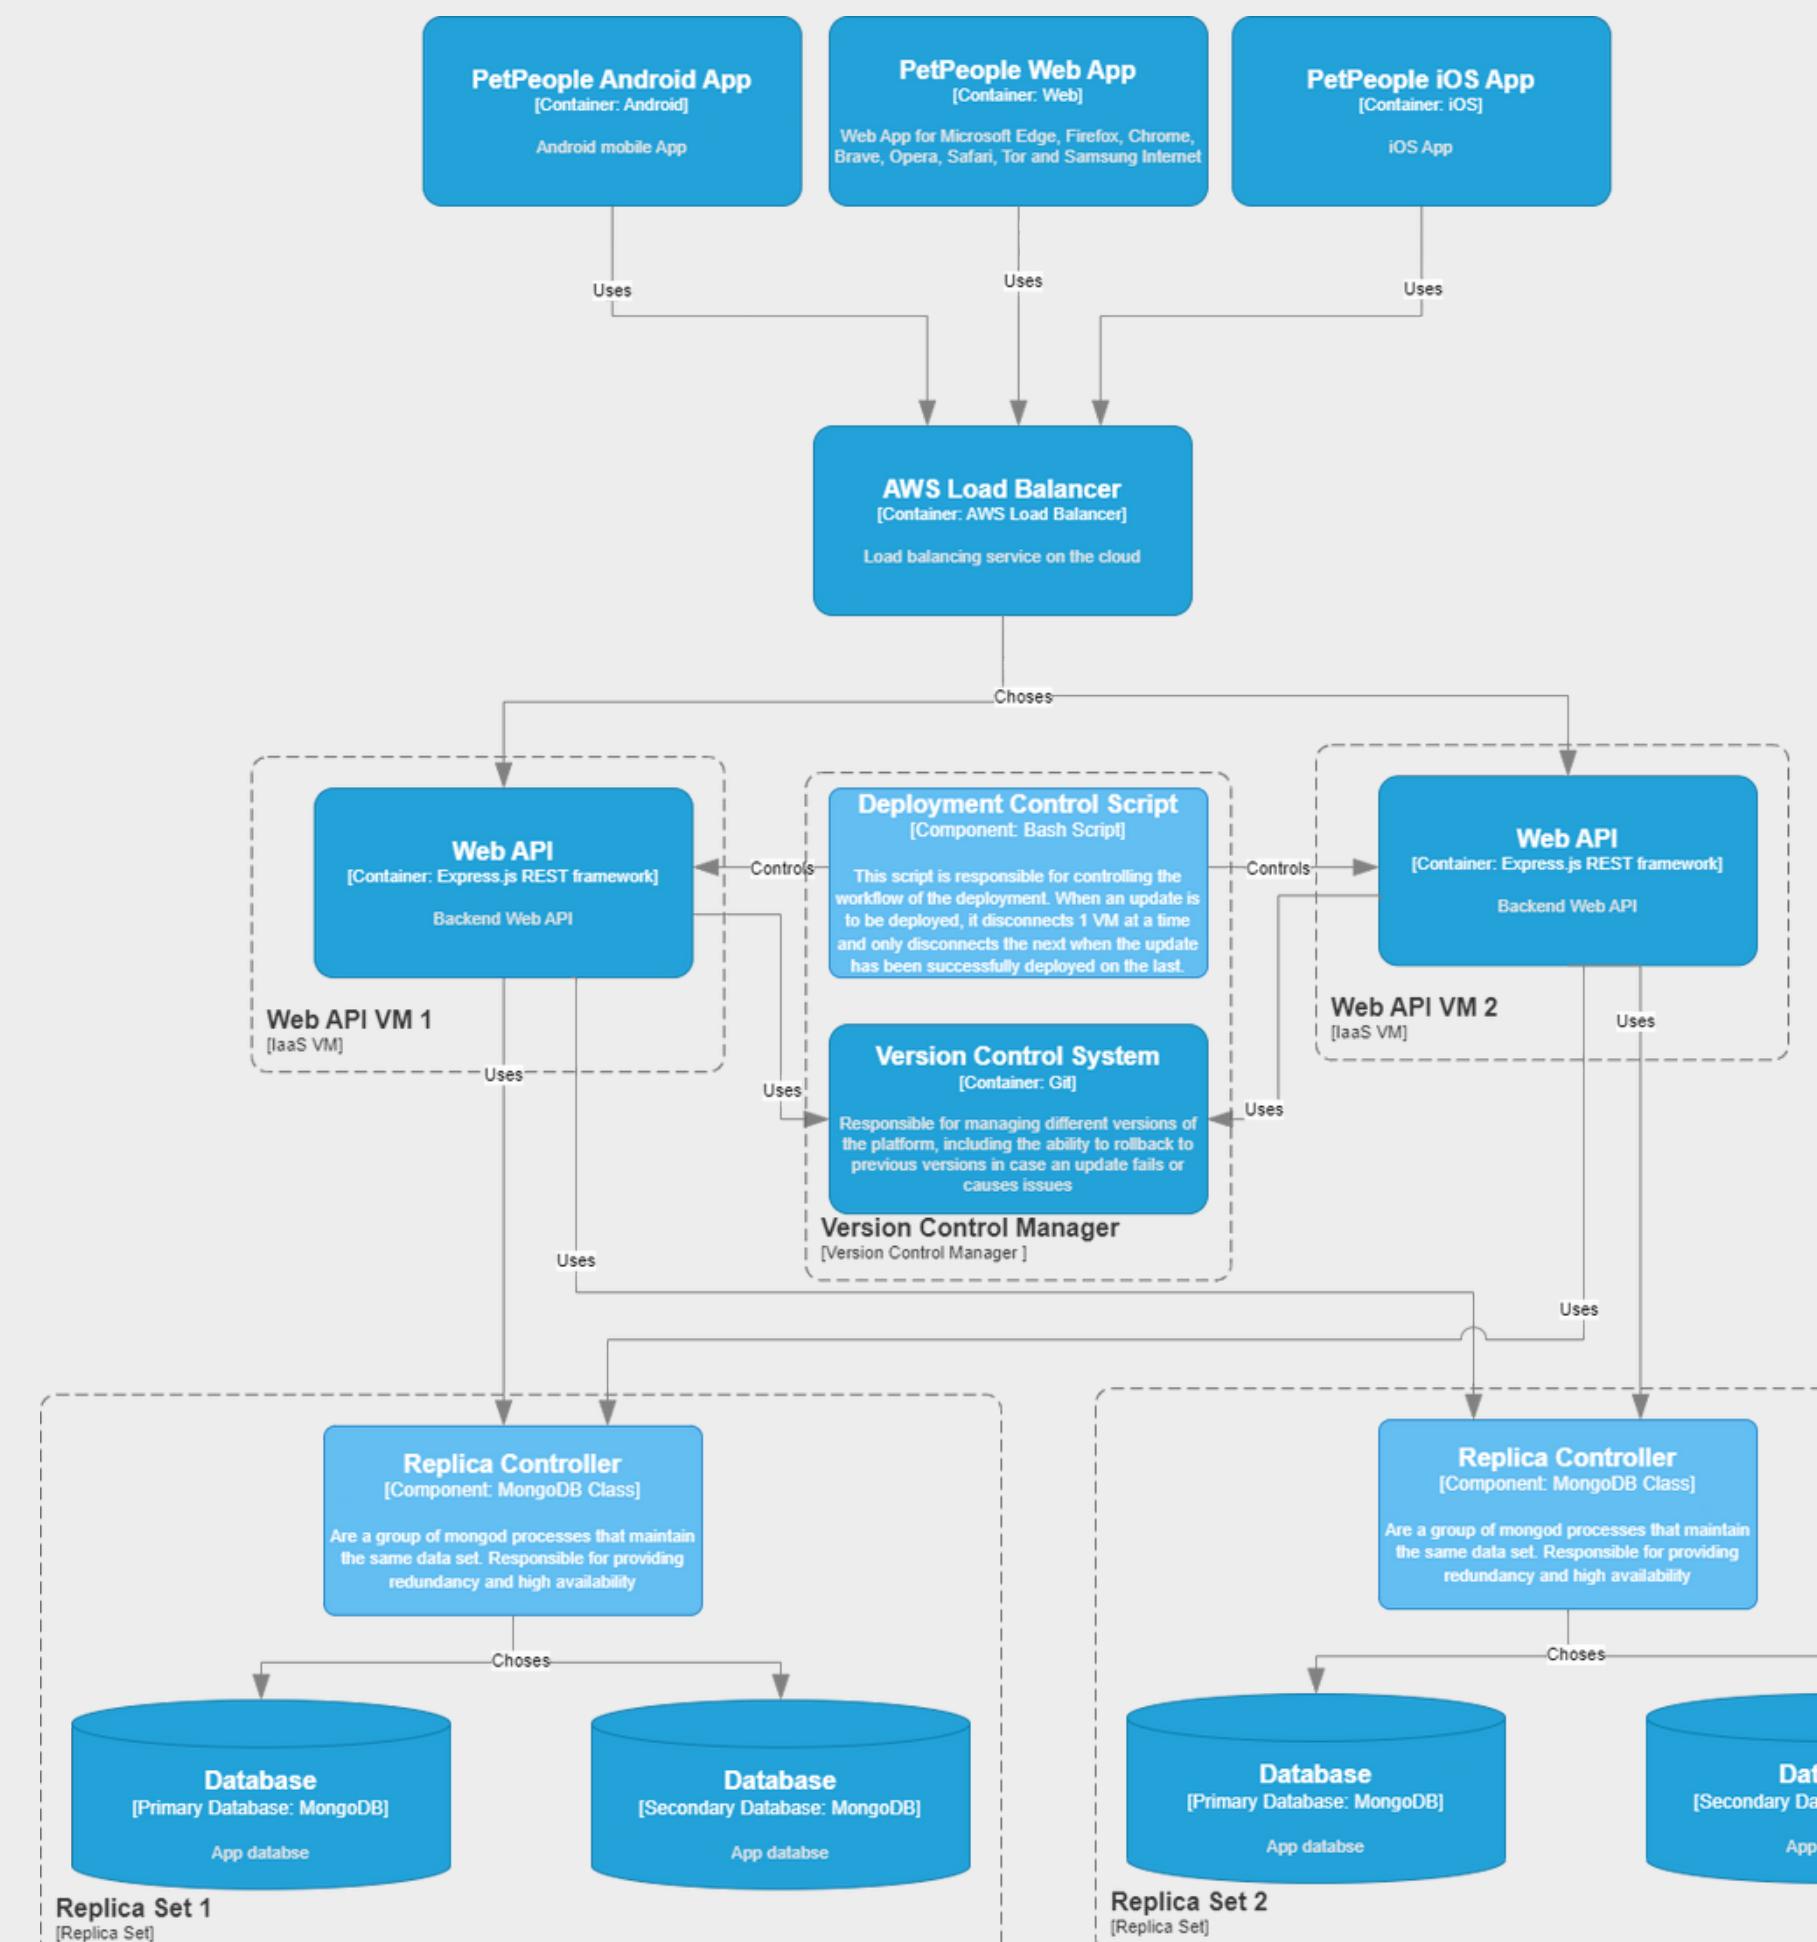
# Scenario



- Up time (Availability)



# Deployment



# Architecture decisions

## Load Balancer



- AWS Load Balancer
  - designed to handle large amounts of traffic;
  - health checks;
  - auto scaling features to help monitor and adjust to which machines to send traffic;
  - integrated in the AWS global network, which has a 99.99% to 99.999% availability per year.
- NGINX
  - designed to handle large amounts of traffic;
  - often used with other tools such as Docker and Kubernetes;
  - lightweight footprint and easily scale.

Both technologies impact on availability is similar, but due to the application being deployed on AWS, it makes a lot more sense to use a AWS load balancer, seeing that it would require some additional architecture components and configurations to use Nginx, which could delay the application's release date impacting BC2-Time to market, which imposes a deadline of 1 year.

# Validation



The advertised uptime of the each of the services, according to documentation

- Load Balancer – 99.99% to 99.999%
- 2 Virtual Machines – 99.99% to 99.999%
- Version Control System (Git repository) – 99.9%
- MongoDB Replicas – 99.995%

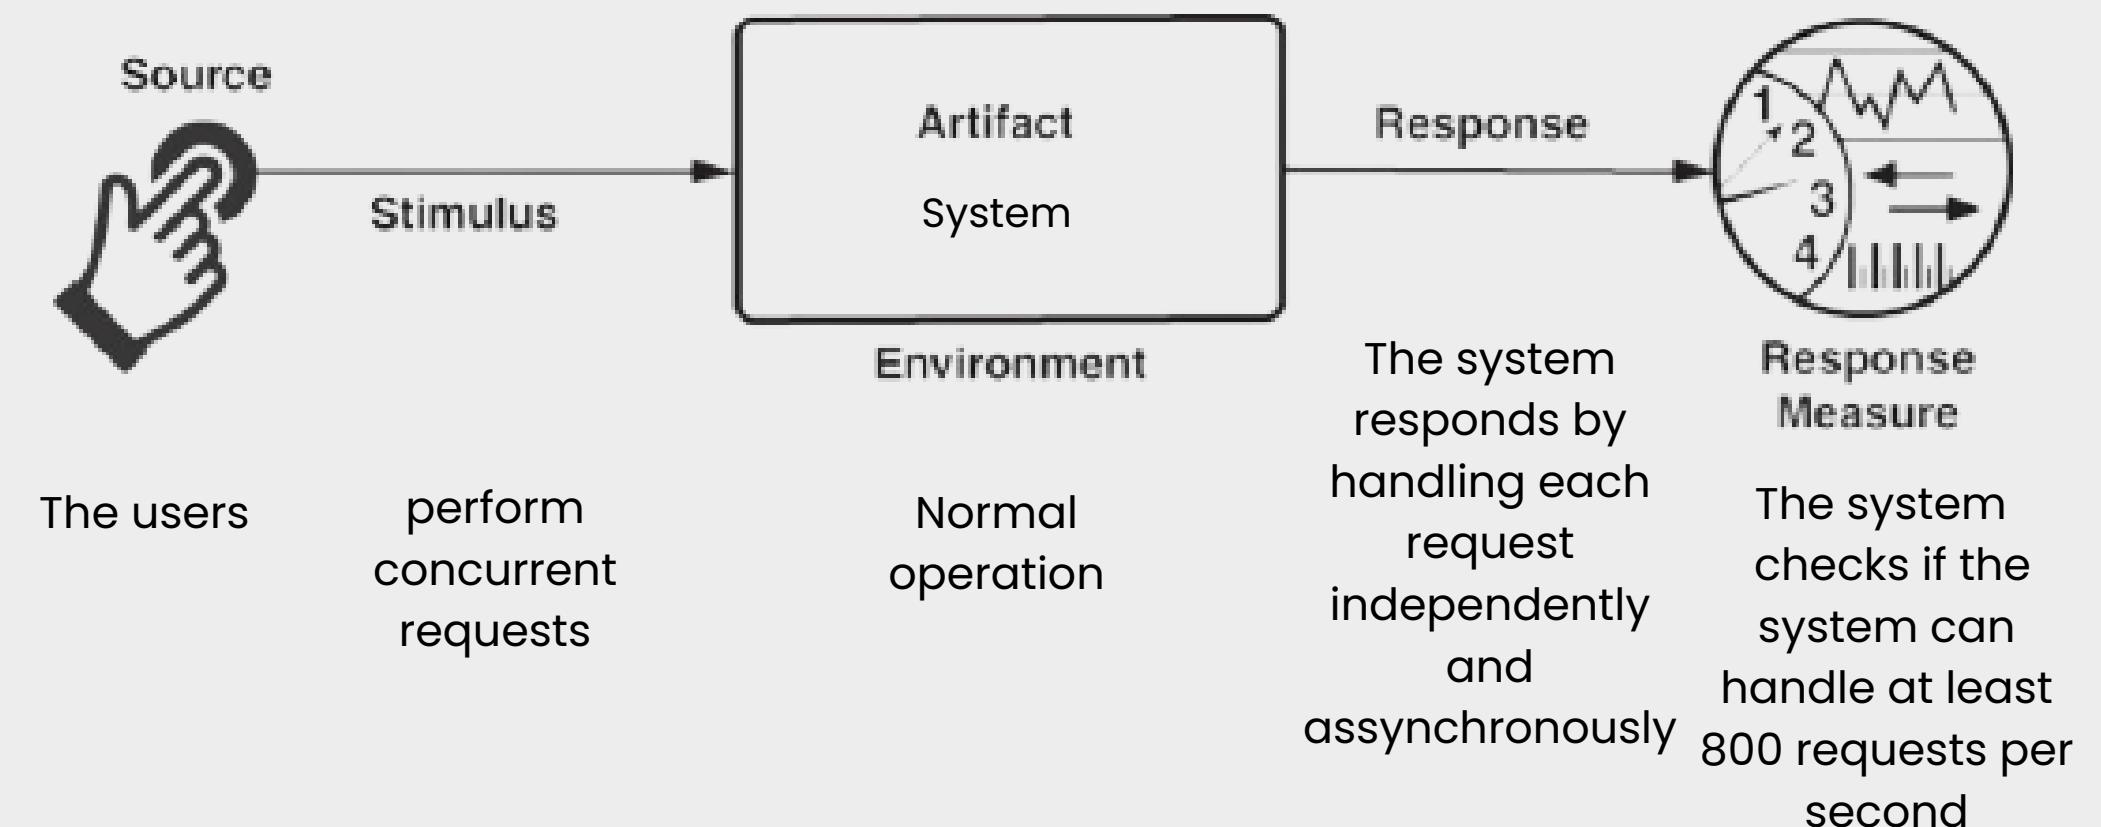
$P((\text{VM1 Down AND Vm2 Down}) \text{ OR GitHub Down OR MongoDB Atlas Down OR Load Balancer})$

	<b>Worse expected downtime</b>	<b>Best expected downtime</b>
<b>Hardware</b>	0.01%	0.001%
<b>Software</b>	0.1342985%	0.10500095%
<b>Availability</b>	99.8557%	99.89399%

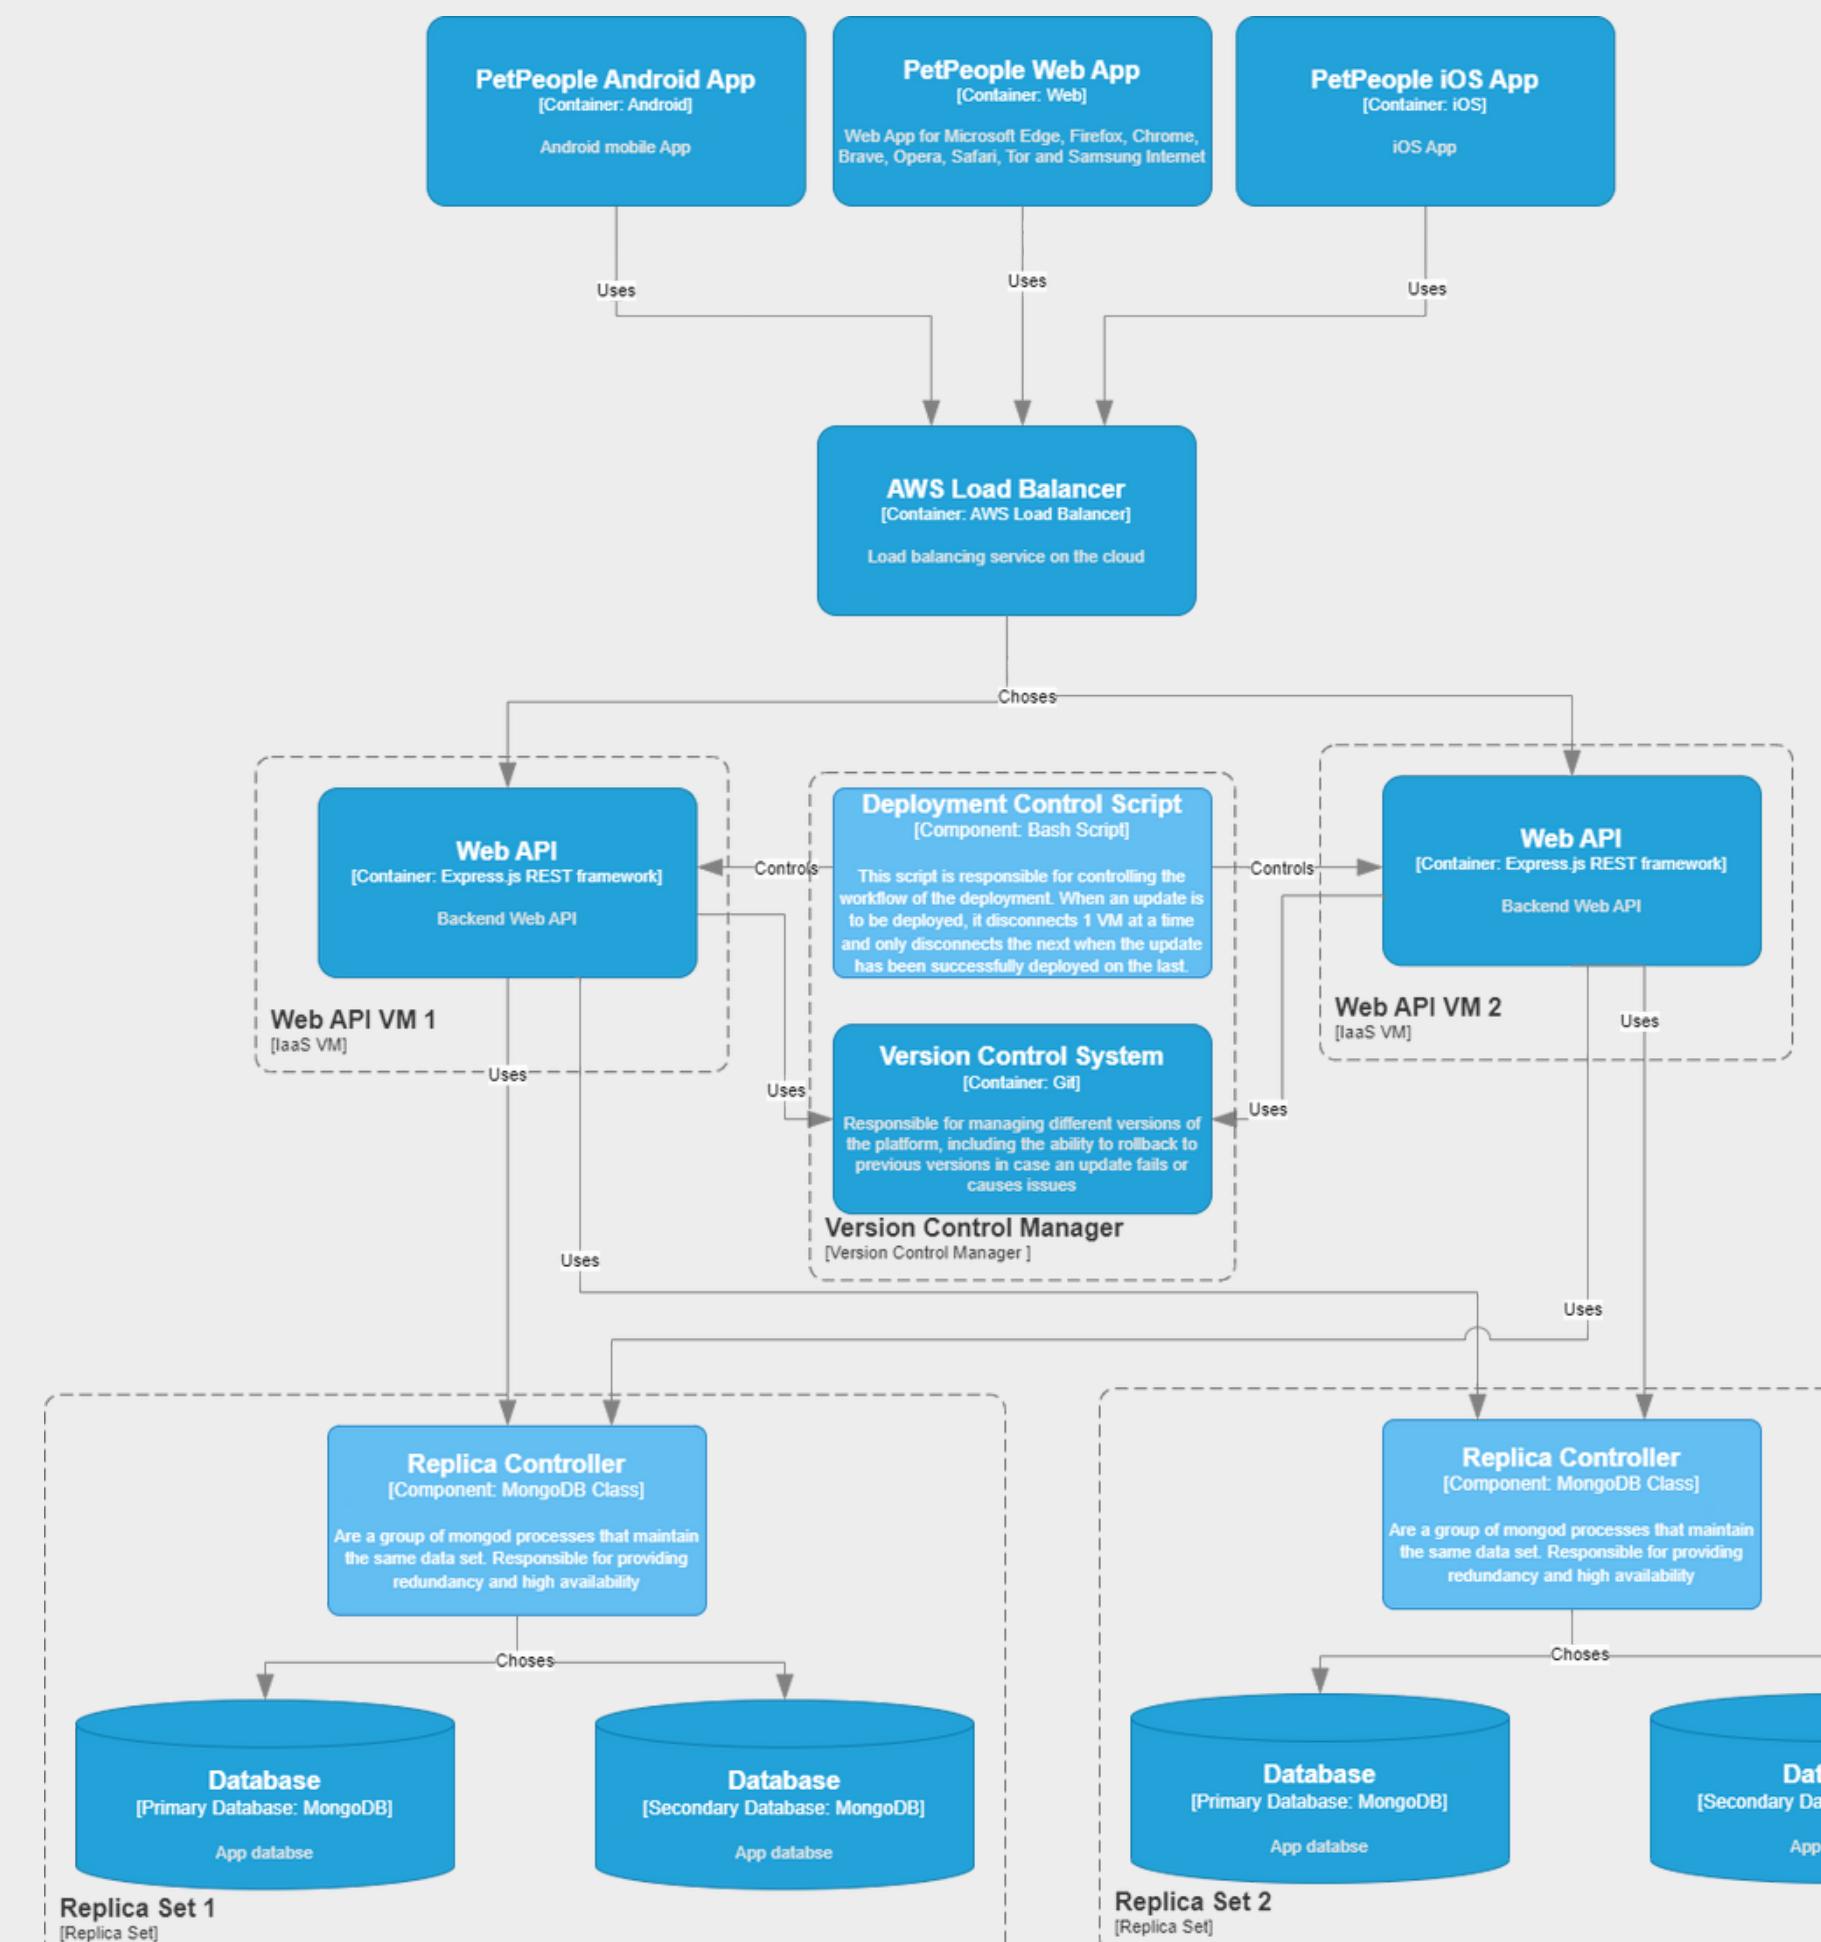
# Scenario



- Number of concurrent requests (Performance)



# Deployment



# Architecture decisions



- Framework
- Load balancer
- Hosting provider
  - Amazon Web Services
  - Google Cloud Platform
  - Microsoft Azure
  - DigitalOcean

Almost all hosting providers offer similar features that guarantee high performance, reliability, security, and scalability. Despite this fact, we believe that Amazon Web Services (AWS) is the most comprehensive hosting provider of them all and offer a wide variety of services in total. Furthermore, we have decided to select additional AWS components; thus our choice was to go with the AWS service.

# Validation

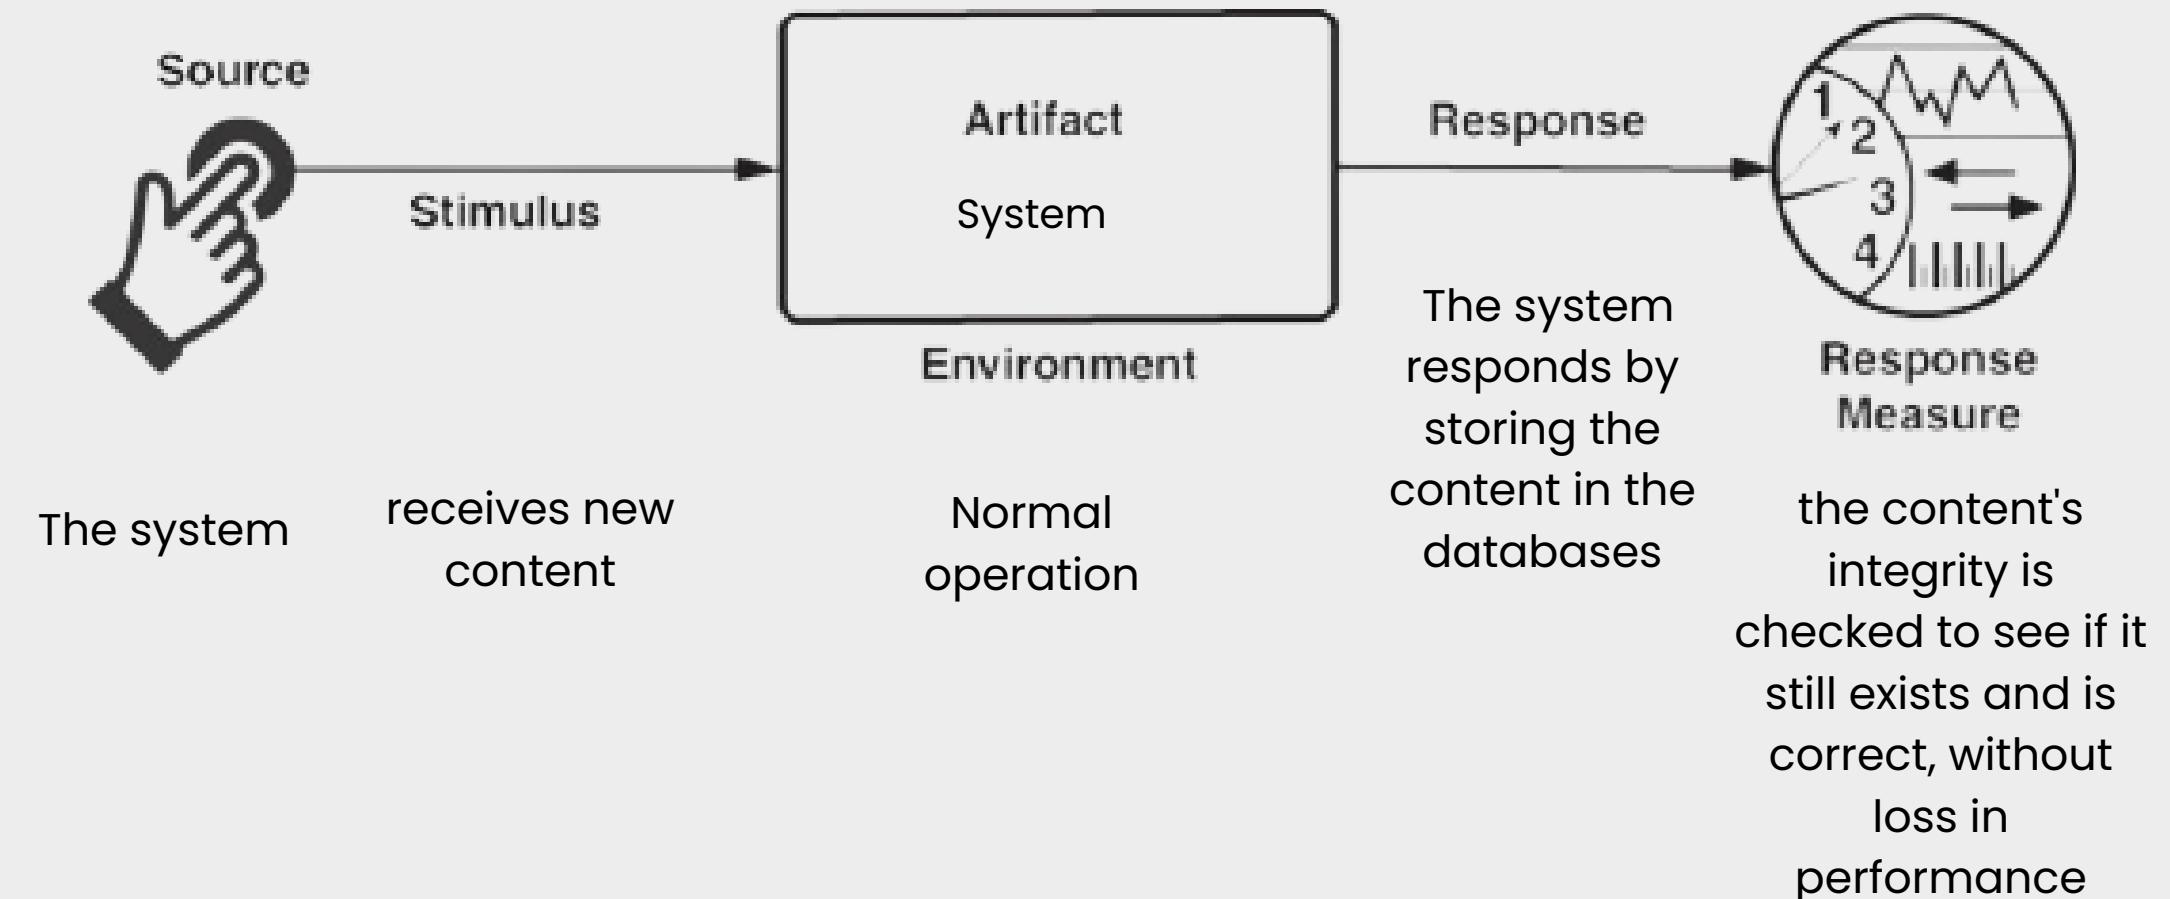


- Amazon EC2 server
  - 4 c6g.4xlarge instances - compute optimized instance that ensures enough power for CPU-intensive workloads
- Elastic Load Balancer - distributes incoming traffic evenly across multiple instances, allowing each instance to handle a portion of the requests and prevent any single instance from becoming overloaded by an abundance of requests
- Amazon Elastic Block Store (EBS)
  - IO2- tailored towards I/O-intensive workloads like databases while providing consistent low-latency performance

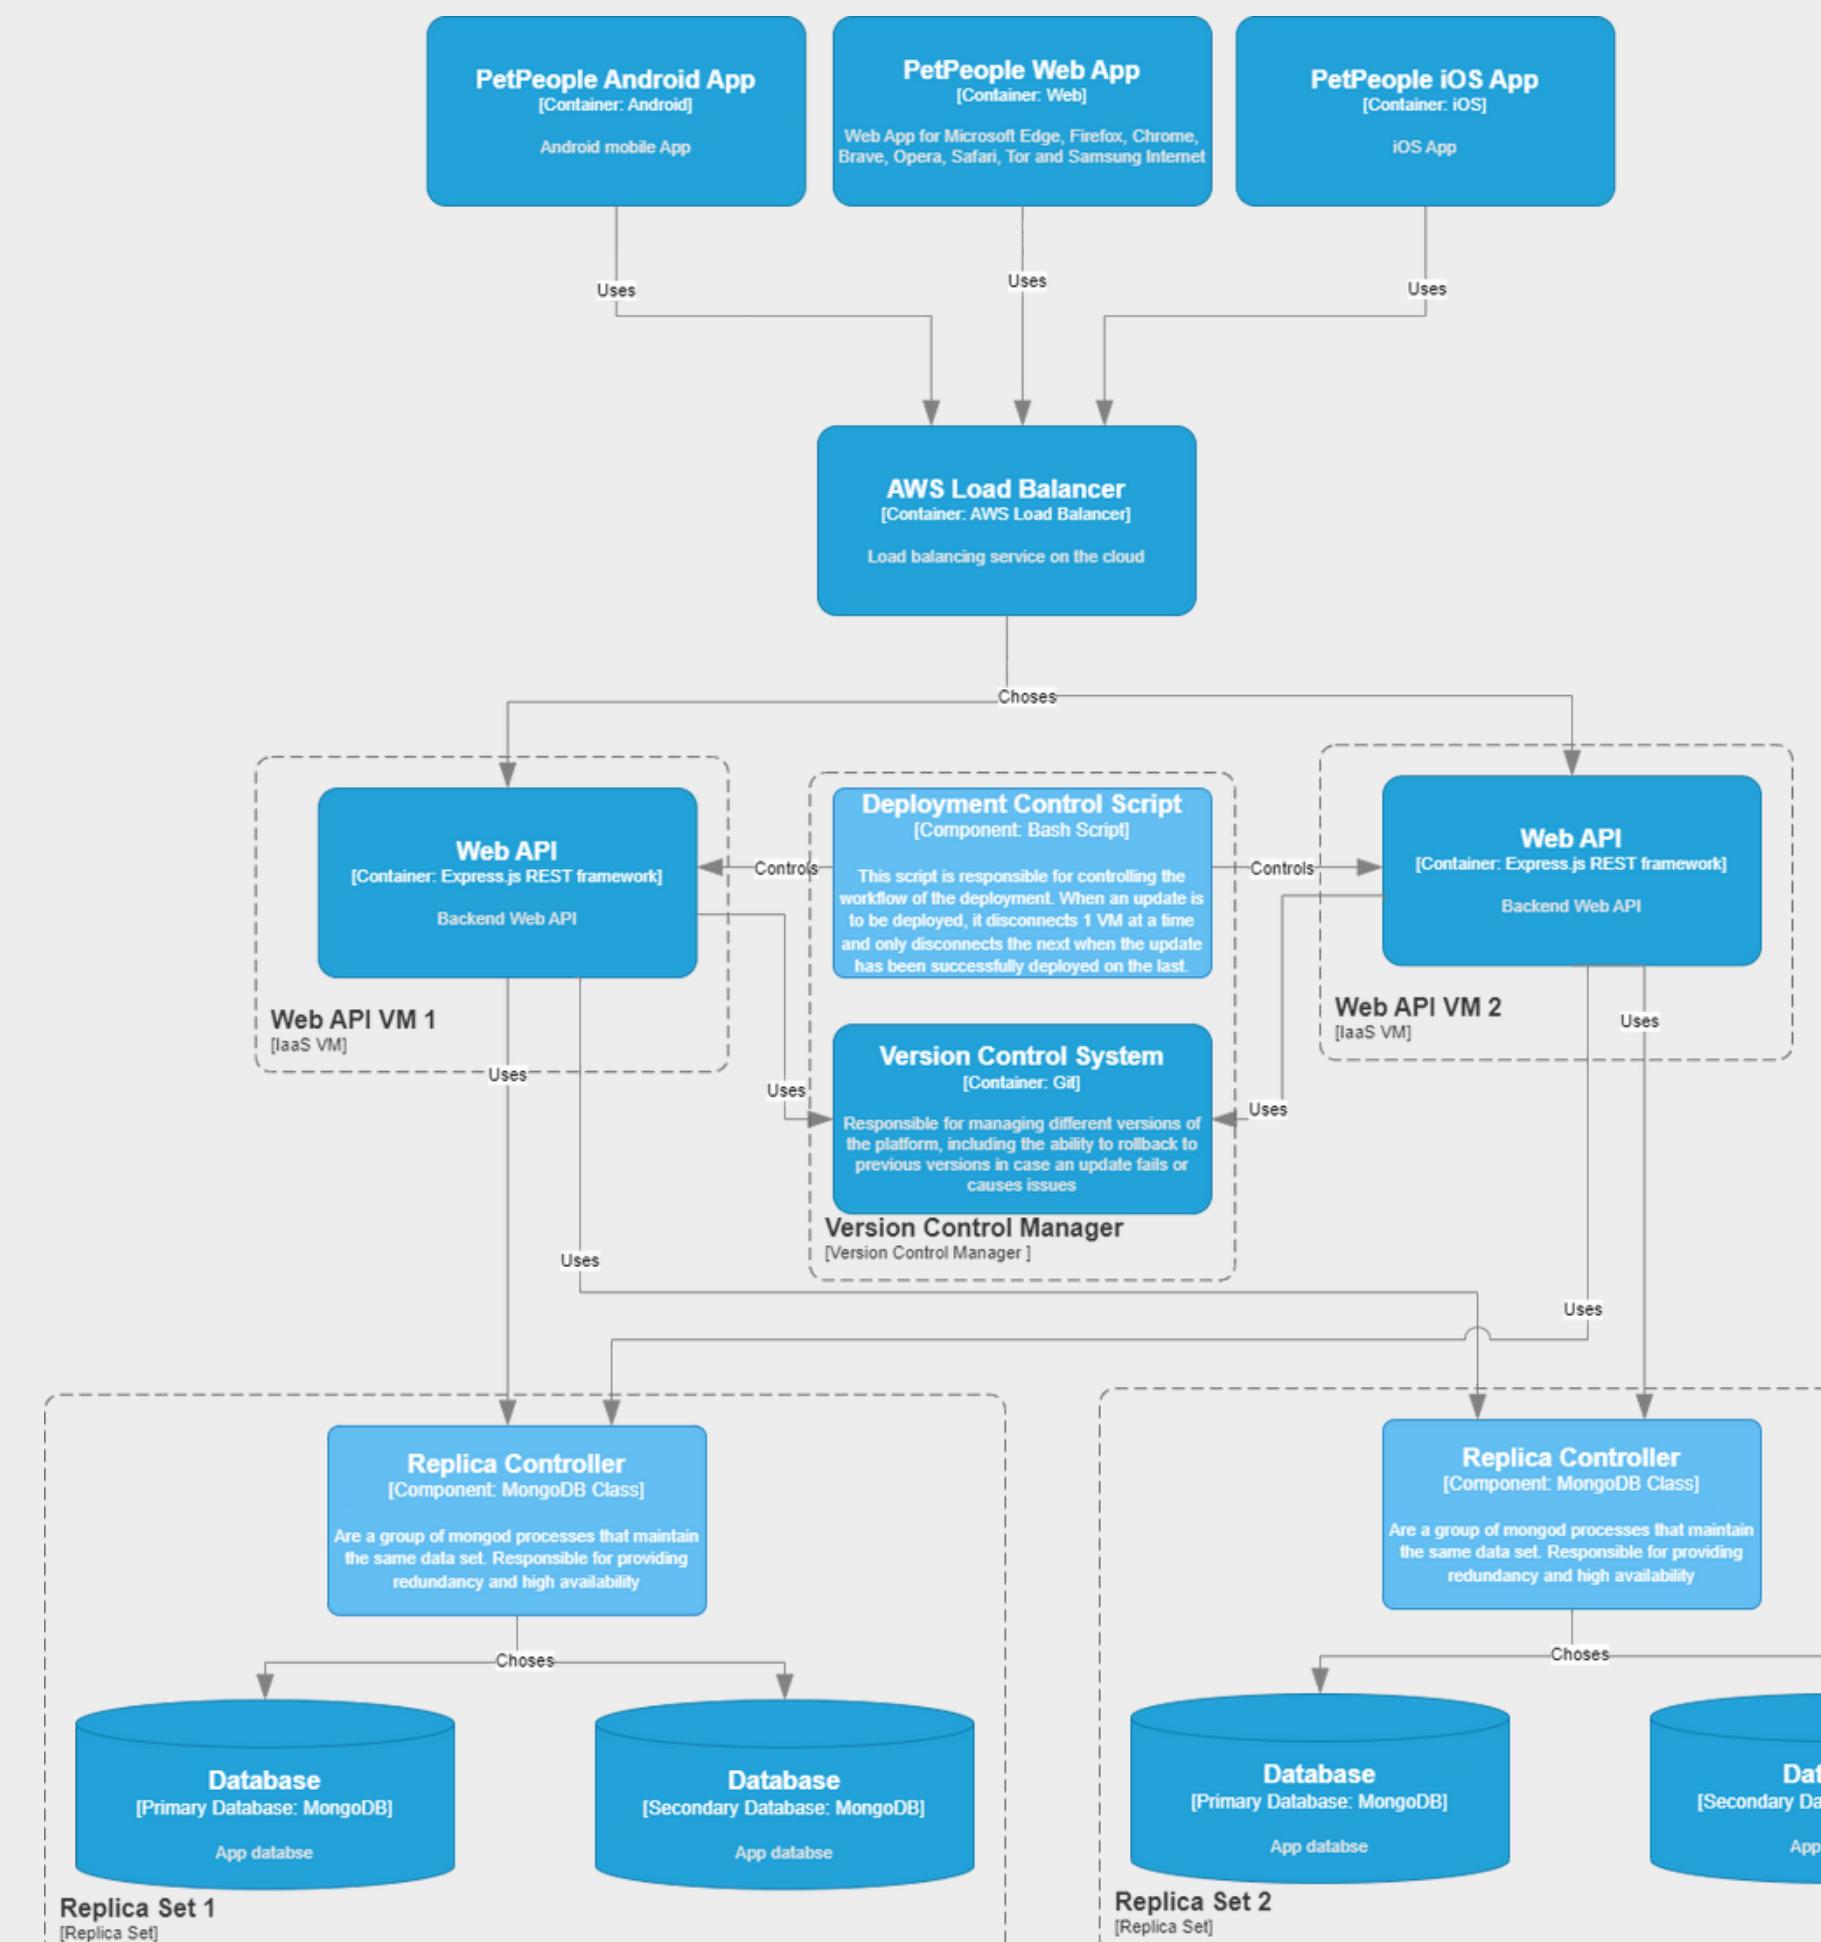
# Scenario



- Sustain an annual content growth of at least 200% in content (Scalability)



# Deployment



# Architecture decisions



- Replica Controller

The replica controller coordinates the behaviour of each replica set and is responsible for electing the current primary database in case of failure of the previous one. It is also responsible for ensuring that the data is synchronized across all the replicas, guaranteeing that any read operation will have the same results, independent of which replica we are currently reading from.

The fact that MongoDB Atlas allows for up to 50 replicas, in multiple replication zones, allows it to scale much more efficiently.

# Validation

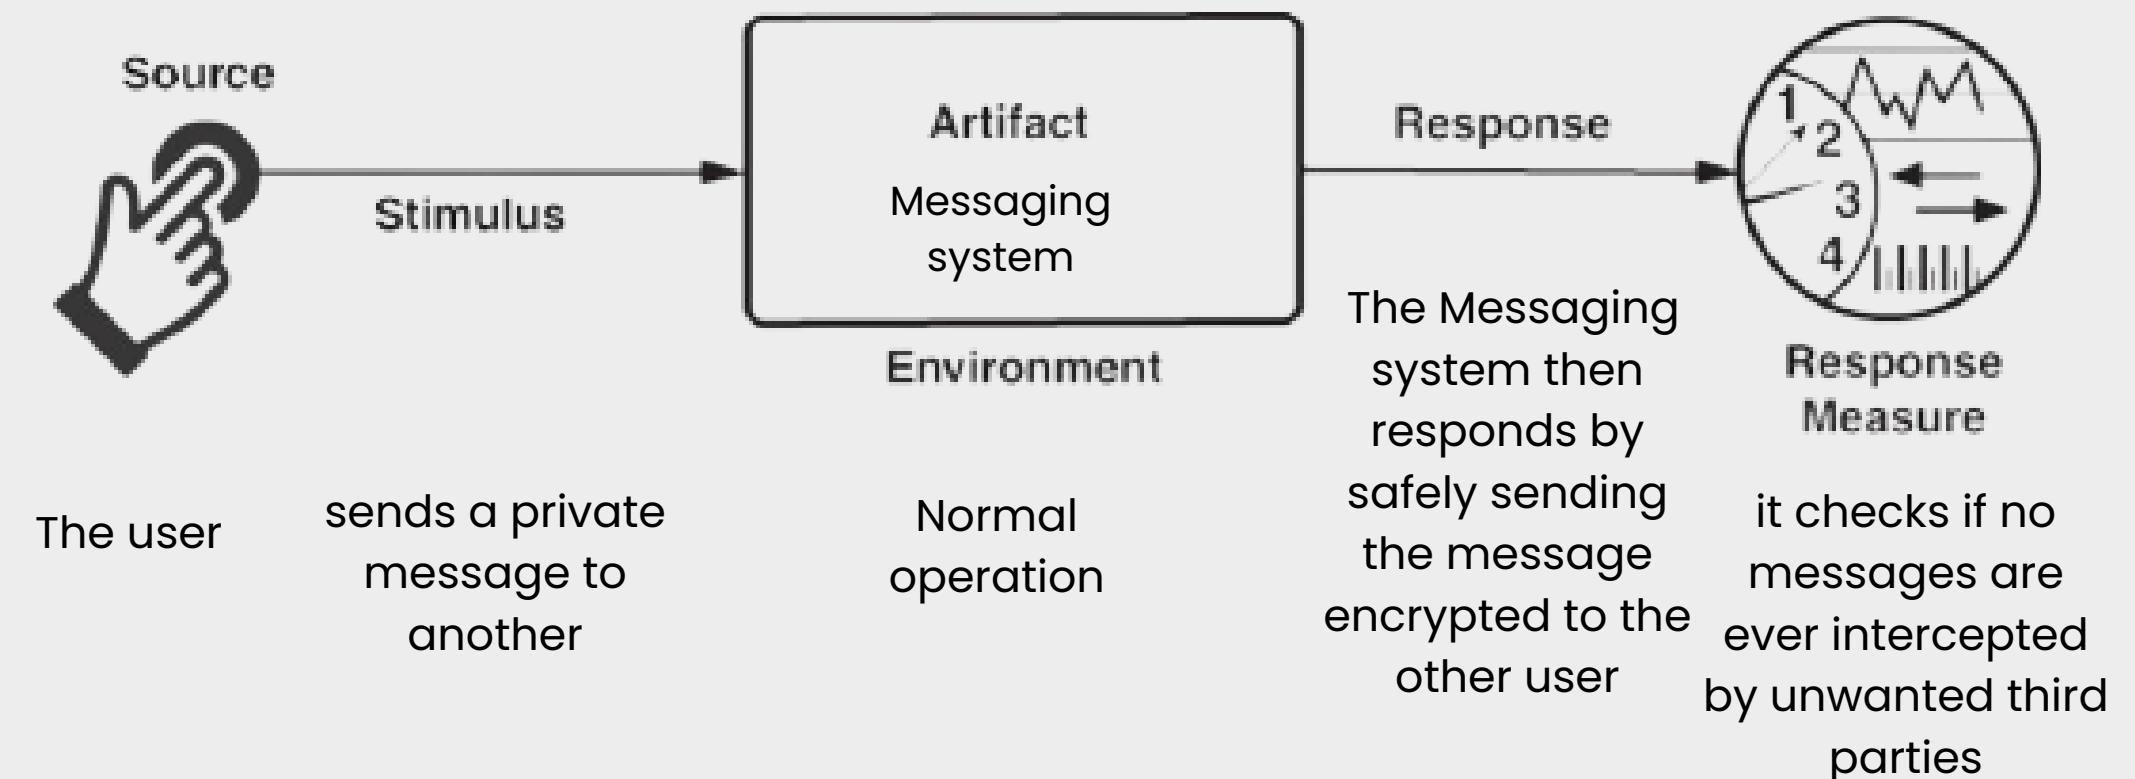


- The amount of data produced in the first year is estimated at 643 terabytes.
- With this in mind, the use of sharding is essential in order to horizontally scale, making it possible to handle large amounts of data.
- According to some documentation, each database can hold up to 64 terabytes of information. We shall use the Amazon Elastic Block Store (EBS) to store the databases, whose maximum storage suffices. The recommended maximum size for a shard is 50 terabytes.
- For our MongoDB cluster, we only need to create new shards when required, having virtually no limit. As such, scalability is assured in our architecture.

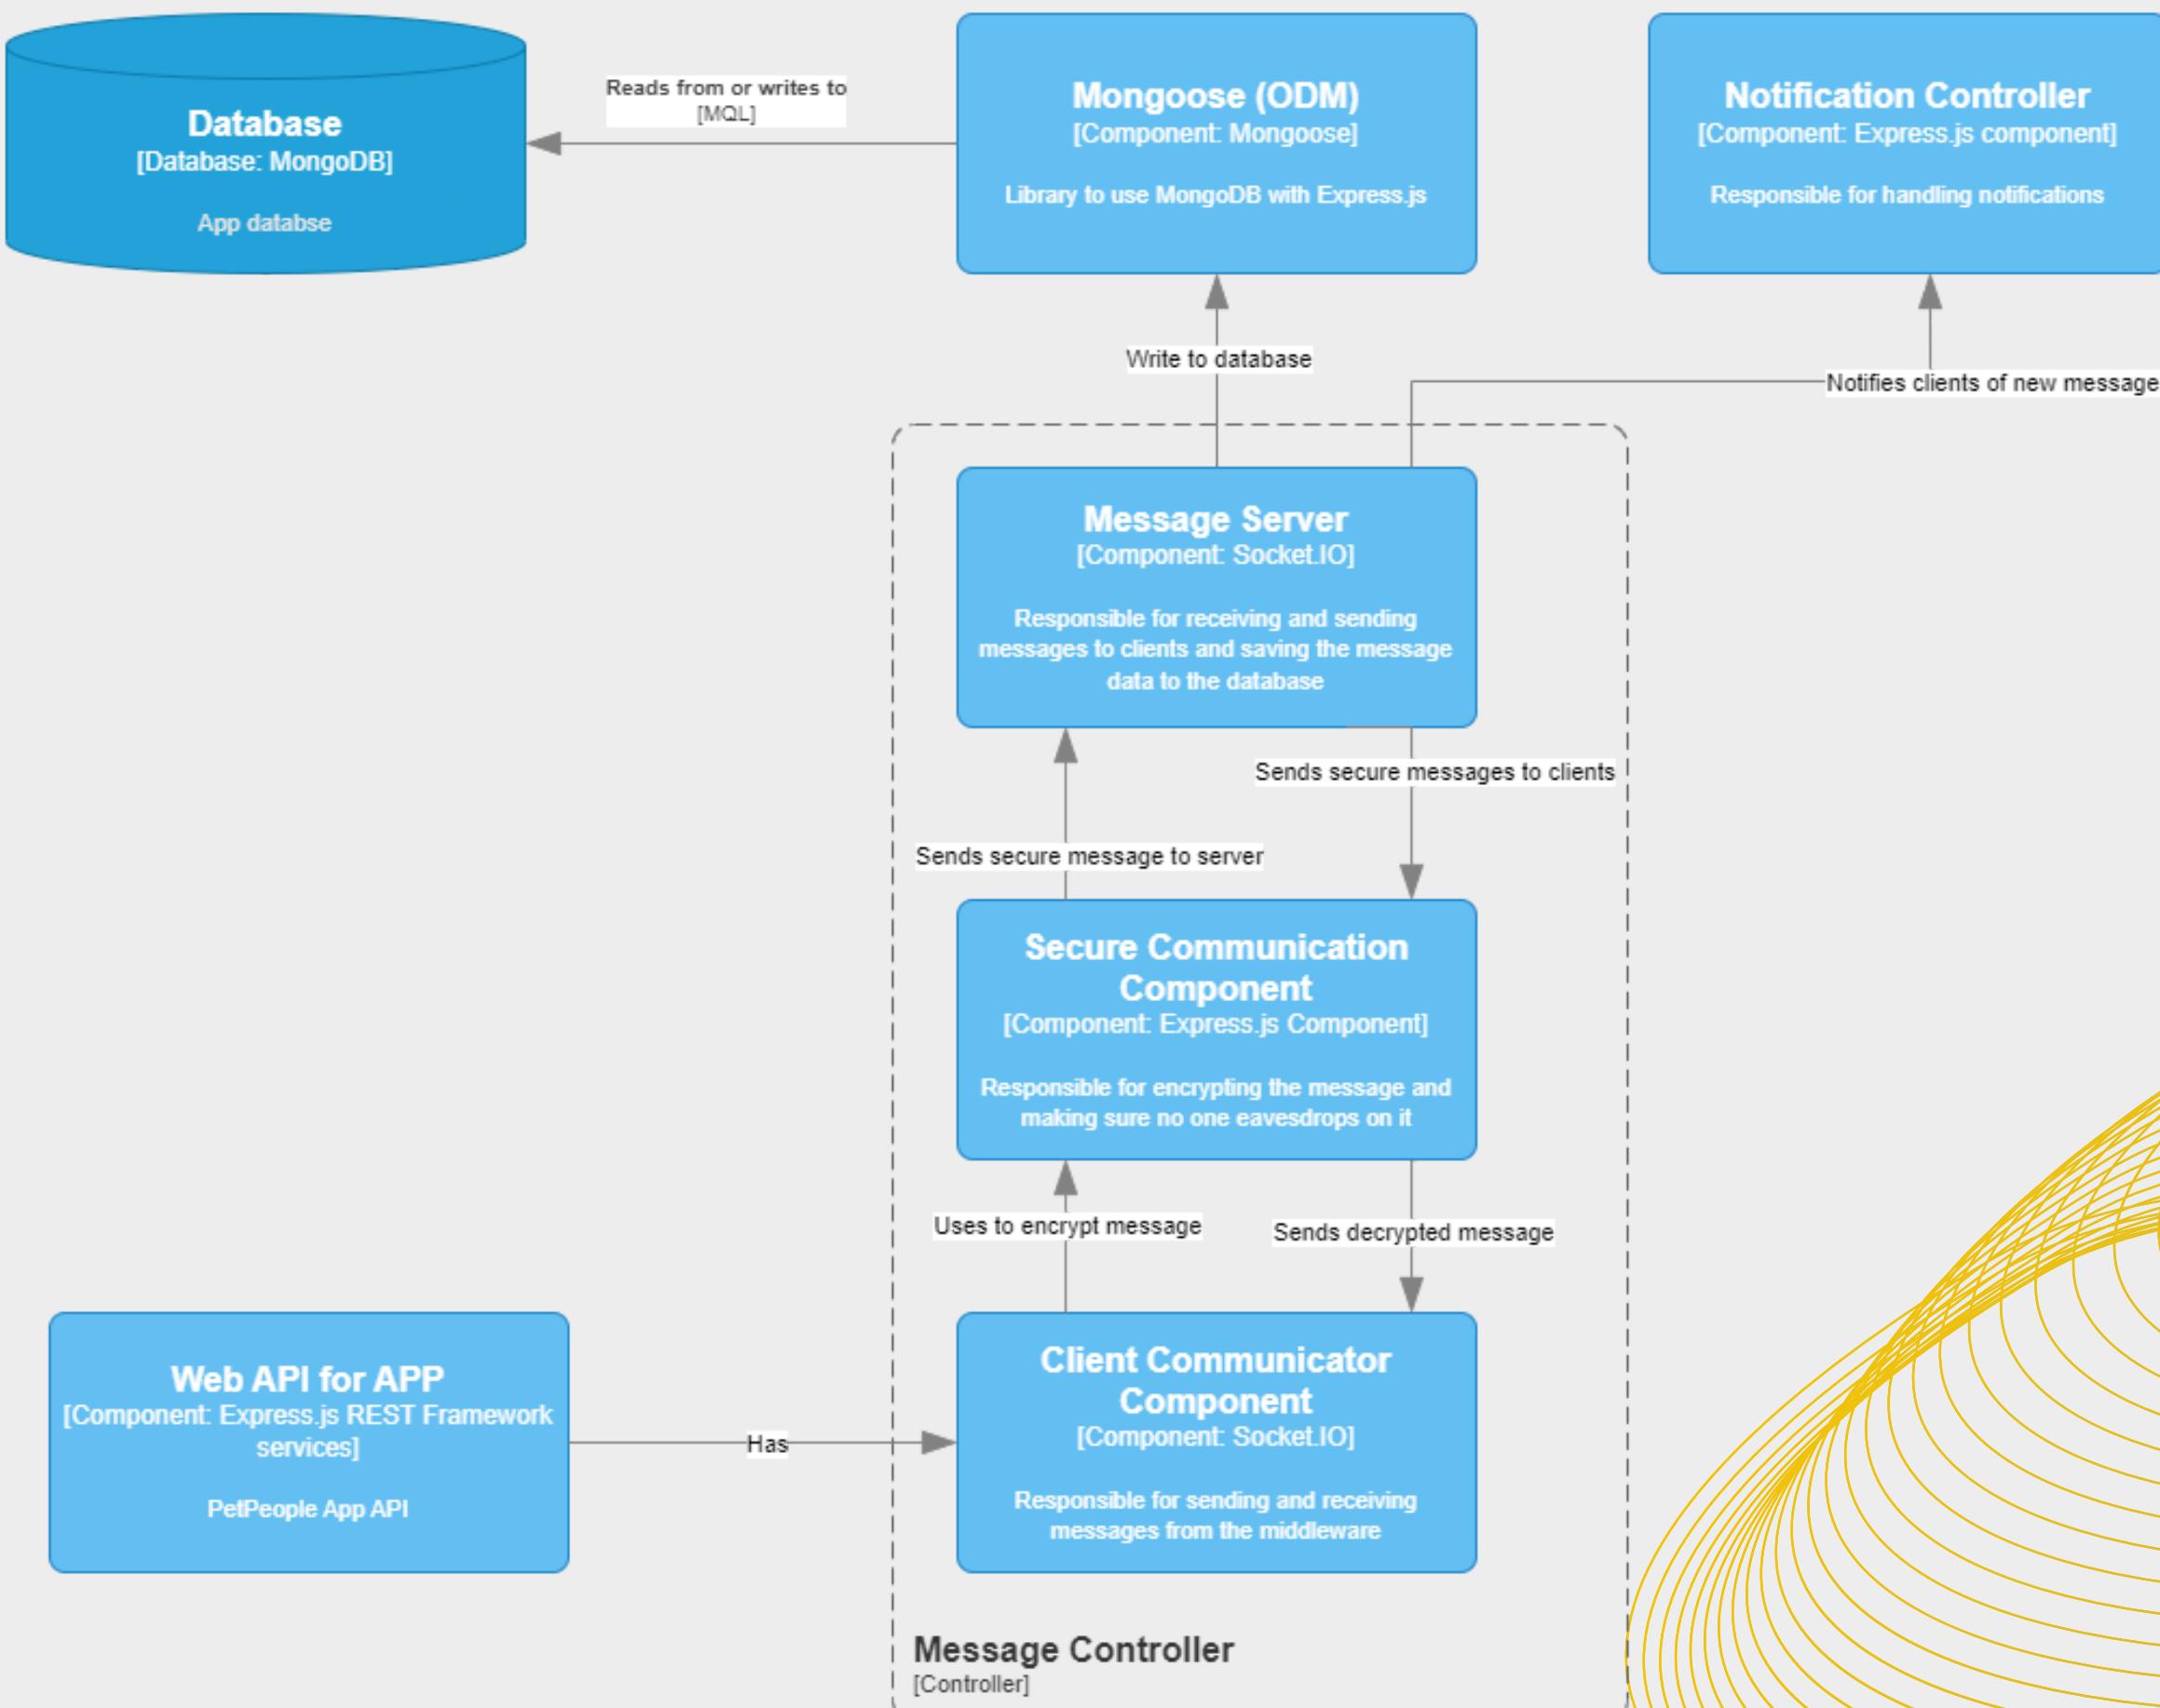
# Scenario



- Secure messaging (Security)



# Message Controller



# Architecture decisions



- End-to-end encryption
  - Is the standard for ensuring a secure and private communications.
  - The data is encrypted on the message sender's device and is only decrypted on the message receiver's device.
  - On top of using this method, the SCC will make use of SSL/TLS encryption protocol to safeguard all transmissions.
- Message-level encryption
  - each message is individually encrypted using a several encryption certificates, making sure that only the receiver can decrypt it using a private key
  - Seeing that it was made having sensitive transaction data in mind (Personal Identification Information, for example) and has a considerable overhead

# Validation



	<b>Security Checklist</b>	<b>Component</b>
User Authentication	1. Verify that users are authenticated before accessing the messaging system	Message Controller
Encryption	4. Verify that messages are encrypted using a secure encryption algorithm	Secure Communication Component
	5. Verify that encrypted messages can only be decrypted by authorized users	Secure Communication Component
	6. Verify that encryption keys are securely stored	Secure Communication Component
Message Integrity	8. Verify that message checksums are used to ensure message integrity	Message Controller
	9. Verify that the messaging system can detect and report any tampering attempts	Secure Communication Component
Access Control	10. Verify that users can only access messages that they are authorized to see	Message Controller

# Risks

<b>Risk Description</b>	<b>Likelihood</b>	<b>Impact</b>	<b>Severity</b>	<b>Mitigation Plan</b>	<b>Owner</b>
Application might not be able to scale to the amount of requests by users. The number of initial MongoDB shards and Amazon Elastic Block Store may not be sufficient	1	5	5	Monitor resource utilization and if it forecasted that there will be a need for more resources, add more shards or increase the size of the Elastic Block	Database Admin
The censor system needs to be tested in the context of animal photos, seeing that Rekognition's documentation is not clear if it detects animal cruelty and abuse. This needs to be verified for a safe and secure platform utilization	3	5	15	Initially have many active admins that respond to reported content. An admin's decision should be informed to the Censor System, so that it can evolve. Once the efficiency of the algorithm has been calculated, the number of admins should be adjusted accordingly. If there are a lot of false positives, then we need to change the Censor System's dataset. If there are a lot of false negatives, then we need more data.	Tester
The architecture may lack streamlined deployment processes, leading to a complex deployment process	5	3	15	Configure CI/CD pipelines to automate the deployment testing process.	Developer
High mobile data usage. Due to the complexity behind the features, it is somewhat likely that the application has a high mobile data usage	3	3	9	Reduce the quality of images and videos in order to reduce the data needed to load them	Developer
There might be performance bottlenecks in the fetching of data. There needs to be verification of the efficiency of the Intelligent Model and its performance in terms of speed.	5	1	5	Have a defined set of preloadable content (may be old posts that are still in cache) for each user that loads in case the Intelligent System is taking too long.	Developer

# Assumptions



- We assume that 1000 ad impressions lead to 8\$ of revenue.

# Dependencies



- External Systems
  - Payment systems
  - Censor System
  - Google Analytics
  - Email system
  - SMS system
  - Notification system

# Process



- 1.Prioritize requirements;
- 2.Choose framework, database and architecture type;
- 3.Follow functional requirements and make model C4;
- 4.Refine the architecture with quality attributes;
- 5.Choose the right test to validate each quality attribute;
- 6.Validate the architecture.

# Team



Edgar Duarte



Sofia Neves



Tatiana Almeida



# Thank you!

*Pet People*